



kubernetes

rkt and Kubernetes

What's new (and coming) with
Container Runtimes and Orchestration



Jonathan Boule

github.com/jonboulle
[@baronboulle](https://twitter.com/baronboulle)



CoreOS

Why rkt and Kubernetes?

~~Why rkt and Kubernetes?~~
Why container runtimes
and orchestration?

CoreOS, Inc (2013 - today)

mission: *"Secure the Internet"*

Started at the OS level: **CoreOS Linux**

- modern, minimal operating system
- self-updating (read-only) image
- updates must be *automatic* and *seamless*

Automatic and seamless

- If the OS is always updating, what about applications running on it?
- Classic use case for *containers* and *orchestration*
 - *containers* decouple the application and OS update lifecycles (update at different cadences)
 - *orchestration* decouples application and OS uptime (services can remain unaffected during OS downtime)

Why container runtimes and orchestration?

So we can provide *seamless updates* and push forward the security of application servers



Why rkt?

The simple answer:
Containers are cool



**A long time ago in an
ecosystem far, far away....**

(2014, to be precise)

2014

- Large incumbent container tool (in CoreOS)
- *Common* practices, but few *best* practices
 - unsigned images (`curl | sudo sh -`)
 - `FROM ubuntu:14.04`
 - PID1 or not to PID1 (zombie reaping problem)
- New platforms emerging, difficult to integrate
 - `systemd + dockerd` = sad times had by all

Enter rkt (and appc)

- Create an alternative container runtime (competition drives innovation)
- Emphasise the importance of *security*
- Spur the conversation around *standards* in the application container ecosystem



a modern, secure container runtime
a simple, composable tool
an implementation of an open standard



a standard application container
open specification
associated tooling



appc

github.com/appc/spec

github.com/appc/acbuild

github.com/appc/docker2aci

github.com/appc/cni

github.com/appc/...



appc

github.com/appc/spec ("appc spec")

github.com/appc/acbuild

github.com/appc/docker2aci

github.com/appc/cni (more on this later..)

[github.com/appc/...](https://github.com/appc/)

appc spec in a nutshell

- Image Format (ACI)
 - what does an application consist of?
- Image Discovery
 - how can an image be located?
- Pods
 - how can applications be grouped and run?
- Executor (runtime)
 - what does the execution environment look like?

appc spec in a nutshell

- Image Format (ACI)
 - what does an application consist of?
- Image Discovery
 - how can an image be located?
- **Pods**
 - **how can applications be grouped and run?**
- Executor (runtime)
 - what does the execution environment look like?

appc pods

- grouping of applications executing in a shared context (network, namespaces, volumes)
- shared fate
- the *only* execution primitive: single applications are modelled as singleton pods

appc pods \approx Kubernetes pods

- grouping of applications executing in a shared context (network, namespaces, volumes)
- shared fate
- the *only* execution primitive: single applications are modelled as singleton pods



a modern, secure container runtime
a simple, composable tool (CLI)
an implementation of an open standard (appc)

rkt - simple CLI tool

no central daemon

no (mandatory) API

apps run directly under spawning process

bash/systemd/kubelet



rkt run ...



application(s)

rkt internals

modular architecture
execution divided into *stages*
stage0 → stage1 → stage2


bash/systemd/kubelet... (invoking process)



rkt (stage0)



pod (stage1)



app1 (stage2)

app2 (stage2)

stage0 (rkt binary)

- primary interface to rkt
- discover, fetch, manage application images
- set up pod filesystems
- manage pod lifecycle
 - rkt run
 - rkt image list
 - rkt gc
 - ...

stage1 (swappable execution engines)

- default implementation
 - based on systemd-nspawn+systemd
 - Linux namespaces + cgroups for isolation
- kvm implementation
 - based on lkvm+systemd
 - hardware virtualisation for isolation
- others?

stage2 (inside the pod)

- actual app execution
- independent filesystems (chroot)
- shared namespaces, volumes, IPC, ...

rkt security ("secure by default")

- image signature verification
- privilege separation
 - e.g. fetch images, expose API (new!) as non-root
- SELinux integration (although </3 overlayfs..)
- lkvm stage1 for true hardware isolation
- TPM attestation (new!)

rkt TPM measurement (new!)

- TPM, Trusted Platform Module: hardware module with cryptographic keys
- Used to "measure" system state
- Historically just use to verify bootloader/OS
- CoreOS added support to GNU Grub
- rkt can now record information about running pods in the TPM

rkt TPM measurement (new!)



rkt API service (new!)

- optional, gRPC-based API daemon
- exposes information on pods and images
- runs as unprivileged user
- read-only
- easier integration with other projects (ahem..)

Why rkt?

Secure
Standards
Composable

Why Kubernetes?

Why Kubernetes?

See the earlier talk...



+



kubernetes

rkt + Kubernetes

- rkt as container runtime (aka "*rktnetes*")
- rkt running Kubernetes ("*rkt fly*")
- rkt networking (CNI)

rkt + Kubernetes

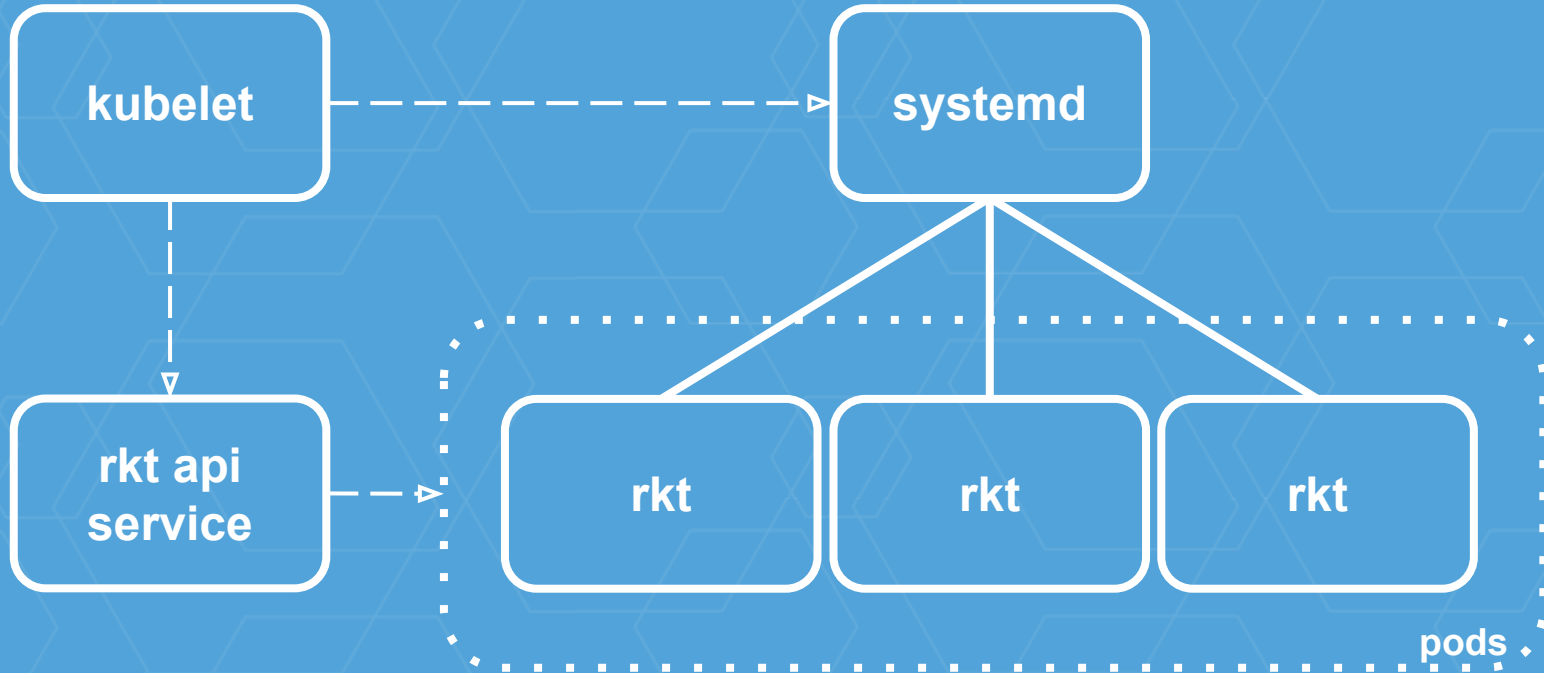
- rkt as container runtime (aka "*rktnetes*")
- rkt running Kubernetes ("*rkt fly*")
- rkt networking (CNI)

rktnetes

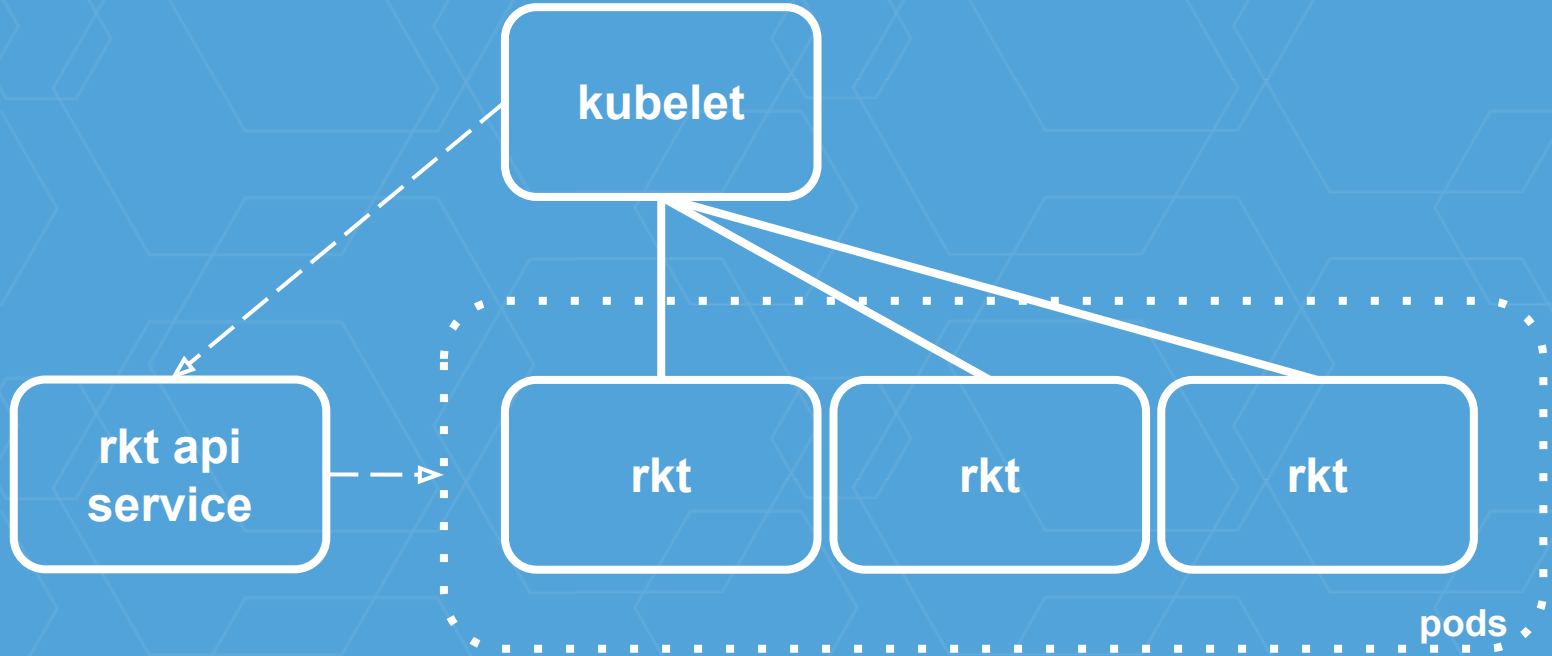
Using rkt as the kubelet's container runtime

- *A pod-native* runtime
- First-class integration with systemd hosts
- *self-contained pods* process model = no SPOF
- Multiple-image compatibility (e.g. docker2aci)
- Transparently swappable container engines

Kubelet + rkt (rktnetes - with systemd)



Kubelet + rkt (rktnetes - without systemd)



Using rkt to run Kubernetes

- Kubernetes components are largely self-hosting, but not entirely
- Need a way to bootstrap kubelet on the host
- On CoreOS, this means in a container..
- ... but kubelet has some unique requirements (like mounting volumes on the host)

Using rkt to run Kubernetes

- rkt "fly" feature (new in rkt 0.15.0)
- unlike `rkt run`, does **not** execute pods
- execute a single application in an unconstrained environment
- all the other advantages of rkt (image discovery, signing/verification, management)

bash/systemd/... (invoking process)

└─▶ rkt (stage0) - without *fly*

└─▶ pod (stage1)

└─▶ app1 (stage2)

└─▶ app2 (stage2)

bash/systemd/... (invoking process)

└─▶ rkt (stage0) - with *fly*

└─▶ application
(kubelet)

rkt networking

Plugin-based
IP(s)-per-pod
Container Networking Interface (CNI)

Container Runtime (e.g. rkt)

Container Networking Interface (CNI)

ptp

macvlan

ipvlan


OVS

CNI: example configuration

```
{  
  "name": "mynet",  
  "type": "ptp",  
  "ipam": {  
    "type": "host-local",  
    "subnet": "10.1.1.0/24"  
  }  
}
```

```
$ rkt run --net=mynet coreos.com/etcd
```


How rkt uses CNI



`/var/lib/rkt/pods/run/$POD_UUID/netns`

Kubernetes networking

Plugin-based (but never left alpha)
IP(s)-per-pod
(sound familiar?)


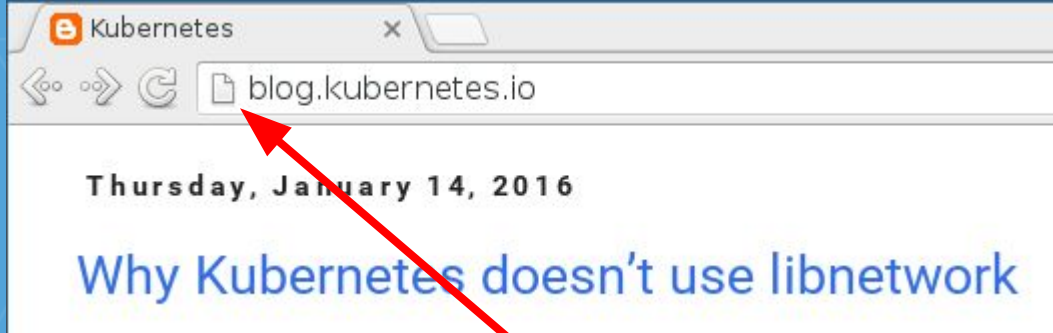
Kubernetes and CNI

Previously CNI was just another plugin type, but soon to be "*the* Kubernetes plugin model"




Kubernetes and CNI

Previously CNI was just another plugin type, but soon to be "*the* Kubernetes plugin model"



Kubernetes and CNI

Previously CNI was just another plugin type, but soon to be "*the* Kubernetes plugin model"



<https://letsencrypt.org/>

Looking ahead

What's coming up for rkt and Kubernetes



First things first...

get things stable

rkt v1.0.0

targeting early February
stable API, CLI, on-disk format
ready to use in production!

rktnetes 1.0

2016Q1

Fully supported, full feature parity
Automated testing on CoreOS


rktnetes 1.0+

LKVM backend by default

Native support for ACIs in Kubernetes

TPM at the Kubernetes level

Tectonic Trusted Computing



<https://coreos.com/blog/coreos-trusted-computing.html>

Kubelet upgrades

- Remember from CoreOS mission:
"updates must be *automatic and seamless*"
- If kubelet is in OS, must be upgraded in lock-step
- But mixed-version clusters don't always work
(e.g. upgrading from 1.07 - 1.1.1: <https://github.com/kubernetes/kubernetes/issues/16961>)

Kubelet upgrades

- Solution: API driven upgrades
- Small agent living on host, invoking kubelet (using rkt fly)
- Reading annotations from the kubelet API server
- Follow along:

<https://github.com/coreos/bugs/issues/1051>

tl;dr:

- Use `rkt` (it's secure, cool, (soon to be) stable)
- Use Kubernetes (for all those earlier reasons)
- Get involved and help define the future of application containers



May 2016 in Berlin

<https://coreos.com/fest> (updated soon!)

- Earlybird tickets
- Sponsorships
- Talk submissions

Questions?

Join us!



github.com/coreos/rkt



Core OS

coreos.com/careers *(now in Berlin!)*