



Overview and Future

Jonathan Boule

CoreOS

@baronboulle



What is etcd?

/etc
distributed

hence, the name...

a clustered **key-value store**

GET and SET operations

a **building block** for higher order systems

primitives for building reliable
distributed systems

What is etcd?

- distributed /etc
 - cluster-level configuration
- clustered key-value store
- primitives for building reliable distributed systems
 - distributed locking system
 - distributed scheduling system
- supports a lot of large distributed applications
 - SkyDNS, Kubernetes, CloudFoundry, ...

History of etcd

History of etcd

- 2013.8 Alpha release (v0.x)
- 2015.2 Stable release (v2.0+)
 - stable replication engine (new Raft implementation)
 - stable v2 API
- 2016.? (v3.0+)
 - efficient, powerful API
 - highly scalable backend

etcd today

- Production-ready!
 - long running failure injection tests
 - no known data loss issues
 - no known inconsistency issues
 - used in critical CoreOS systems like locksmith and fleet
 - trusted by Google, Pivotal, compose and many more

Why build etcd?

Why build etcd?

- CoreOS mission: "Secure the internet"
- Updating servers = rebooting servers
- Move towards new application container paradigm
- Need a:
 - shared configuration store (for service discovery etc)
 - distributed lock manager (to co-ordinate reboots)
- Existing solutions were inflexible (undocumented binary API), difficult to configure



Why use etcd?

Why use etcd?

- Highly available
- Highly reliable
- Strong consistency guarantees
- Simple, fast HTTP API
- Open source

"For the most critical data of a distributed system"


**How does etcd
work?**

How does etcd work?

- **Raft** consensus algorithm
 - Using a replicated log to model a state machine
 - *"In Search of an Understandable Consensus Algorithm"* (Ongaro, 2014)
- Three key concepts
 - *Leaders*
 - *Elections*
 - *Terms*

How does etcd work?

- The cluster *elects* a leader for every given *term*
- All log appends (--> state machine changes) are decided by that leader and propagated to followers
- Much much more at <http://raft.github.io/>



How does etcd work?

- Written in Go, statically linked
- `/bin/etcd`
 - daemon
 - 2379 (client requests/HTTP + JSON API)
 - 2380 (peer-to-peer/HTTP + protobuf)
- `/bin/etcdctl`
 - command line client
- `net/http`, `encoding/json`, `golang/protobuf`, ...




etcd basics

clusters

Available



Available



Available



Unavailable



etcd cluster failure tolerance

| CLUSTER SIZE | MAJORITY | FAILURE TOLERANCE |
|--------------|----------|-------------------|
| 1 | 1 | 0 |
| 3 | 2 | 1 |
| 4 | 3 | 1 |
| 5 | 3 | 2 |
| 6 | 4 | 2 |
| 7 | 4 | 3 |
| 8 | 5 | 3 |
| 9 | 5 | 4 |

etcd basics

API

Simple HTTP API (v2)

- GET `/v2/keys/foo`
 - Get the value of a key
- GET `/v2/keys/foo?wait=true`
 - Wait for changes on key foo
 - Receive notifications on every change
- PUT `/v2/keys/foo -d value=bar`
 - Set the value of a key
- DELETE `/v2/keys/foo`
 - Delete a key

Compare-and-Swap

PUT /v2/keys/foo?prevValue=bar -d
value=ok

```
CAS(/foo, bar, ok)
```

```
if /foo == bar
```

```
    set(/foo, ok)
```

```
else
```

```
    do nothing
```

Compare-and-Delete

DELETE /v2/keys/foo?prevValue=bar

```
CAD(/foo, bar)
```

```
if /foo == bar
```

```
    delete(/foo)
```

```
else
```

```
    do nothing
```

Simple HTTP API (v2)

Native Go bindings

```
import "github.com/coreos/etcd/client"

cl := client.New(client.Config{})
kapi := client.NewKeysAPI(cl)
kapi.Set("foo", "bar", ...)
```



etcd apps



etcd apps
locksmith

locksmith

- cluster wide reboot lock
 - "semaphore for reboots"
- CoreOS updates happen *automatically*
 - prevent all the machines restarting at once...

Cluster Wide Reboot Lock



server1

server2

server3

Cluster Wide Reboot Lock



server1

server2

server3

needs reboot

Cluster Wide Reboot Lock

- Need to reboot? Decrement the semaphore key (*atomically*) with etcd
- `manager.Reboot()` and wait...
- After reboot, increment the semaphore key in etcd (*atomically*)


Cluster Wide Reboot Lock

Sem=1




needs reboot

Cluster Wide Reboot Lock





Cluster Wide Reboot Lock



Cluster Wide Reboot Lock




Sem=0



Reboot()

Cluster Wide Reboot Lock


Sem=0




Rebooting...

Cluster Wide Reboot Lock

Sem=0




needs reboot




Rebooting...

Cluster Wide Reboot Lock



Cluster Wide Reboot Lock




Cluster Wide Reboot Lock



Sem=0



wait...



needs reboot




Rebooting...

Cluster Wide Reboot Lock



Sem=0



wait...




needs reboot




Rebooted!


Cluster Wide Reboot Lock




Cluster Wide Reboot Lock



Cluster Wide Reboot Lock






Cluster Wide Reboot Lock



Cluster Wide Reboot Lock

Sem=0



Reboot()



etcd apps
flannel

flannel

- virtual (overlay) network, providing a subnet to each host, and handling all routing
- for use with Kubernetes etc
- uses etcd to store network configuration, allocated subnets, and auxiliary data (host's IP).



etcd apps
skydns

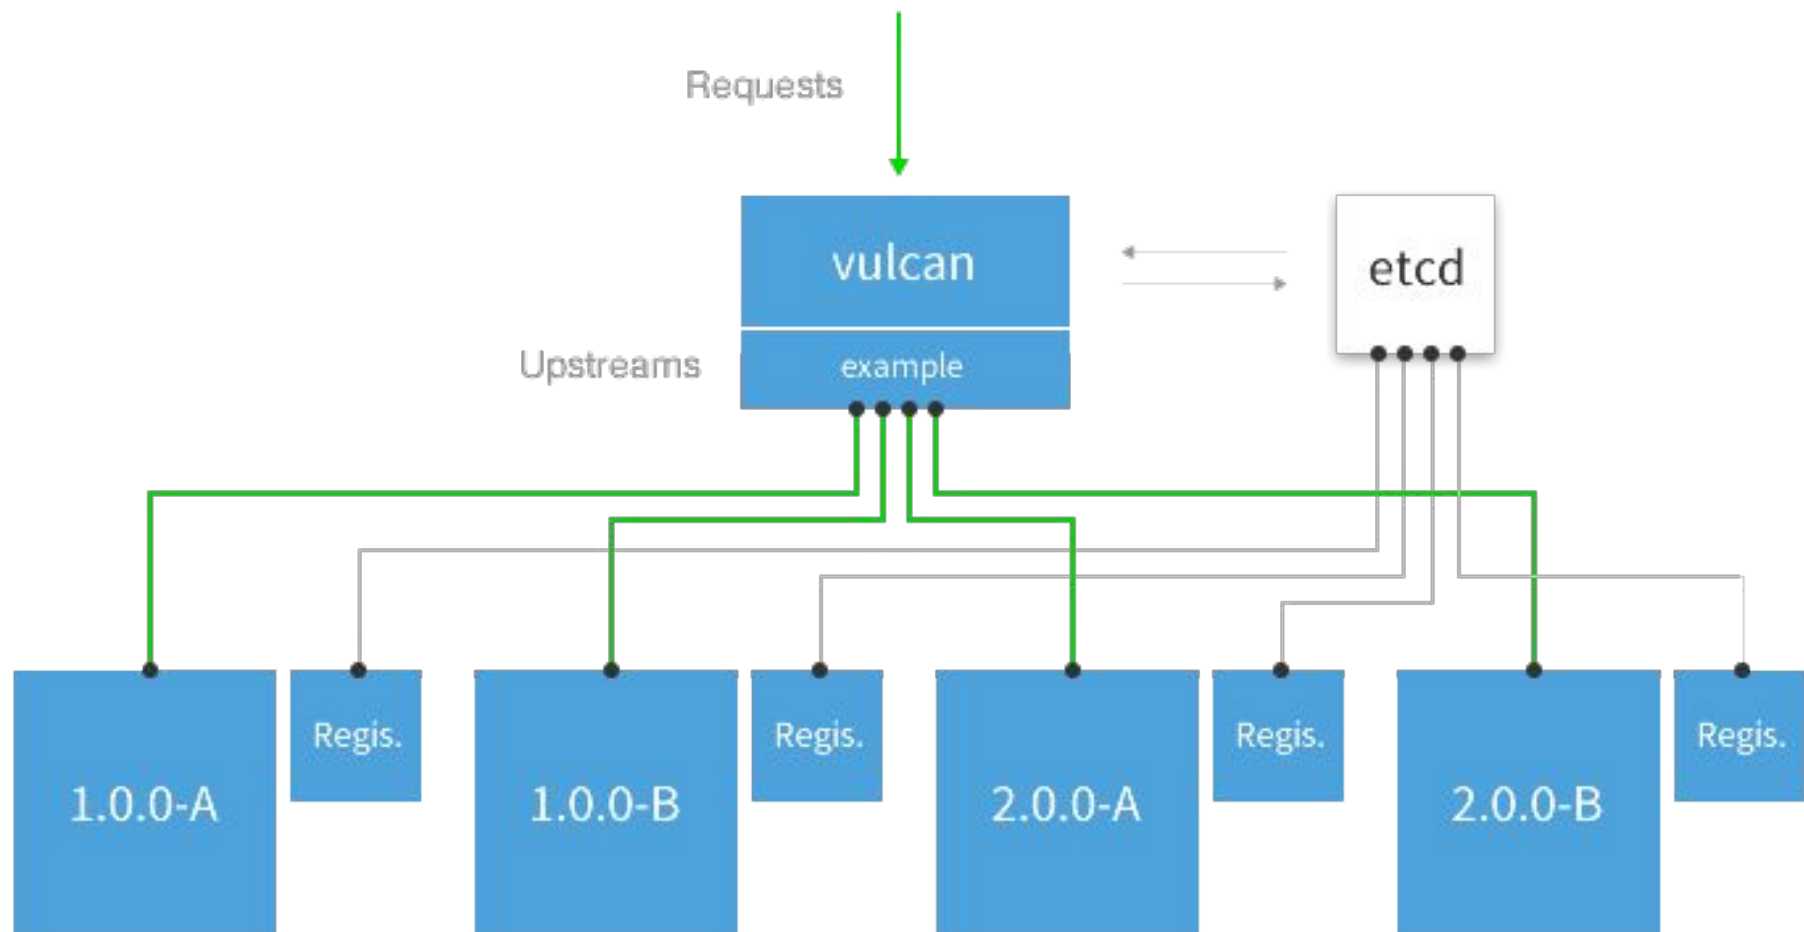
skydns

- Service discovery and DNS server
- backed by etcd for all configuration/records

etcd apps
vulcand

vulcand

- "programmable, extendable proxy for microservices"
- HTTP load balancer
- etcd for storing all configuration





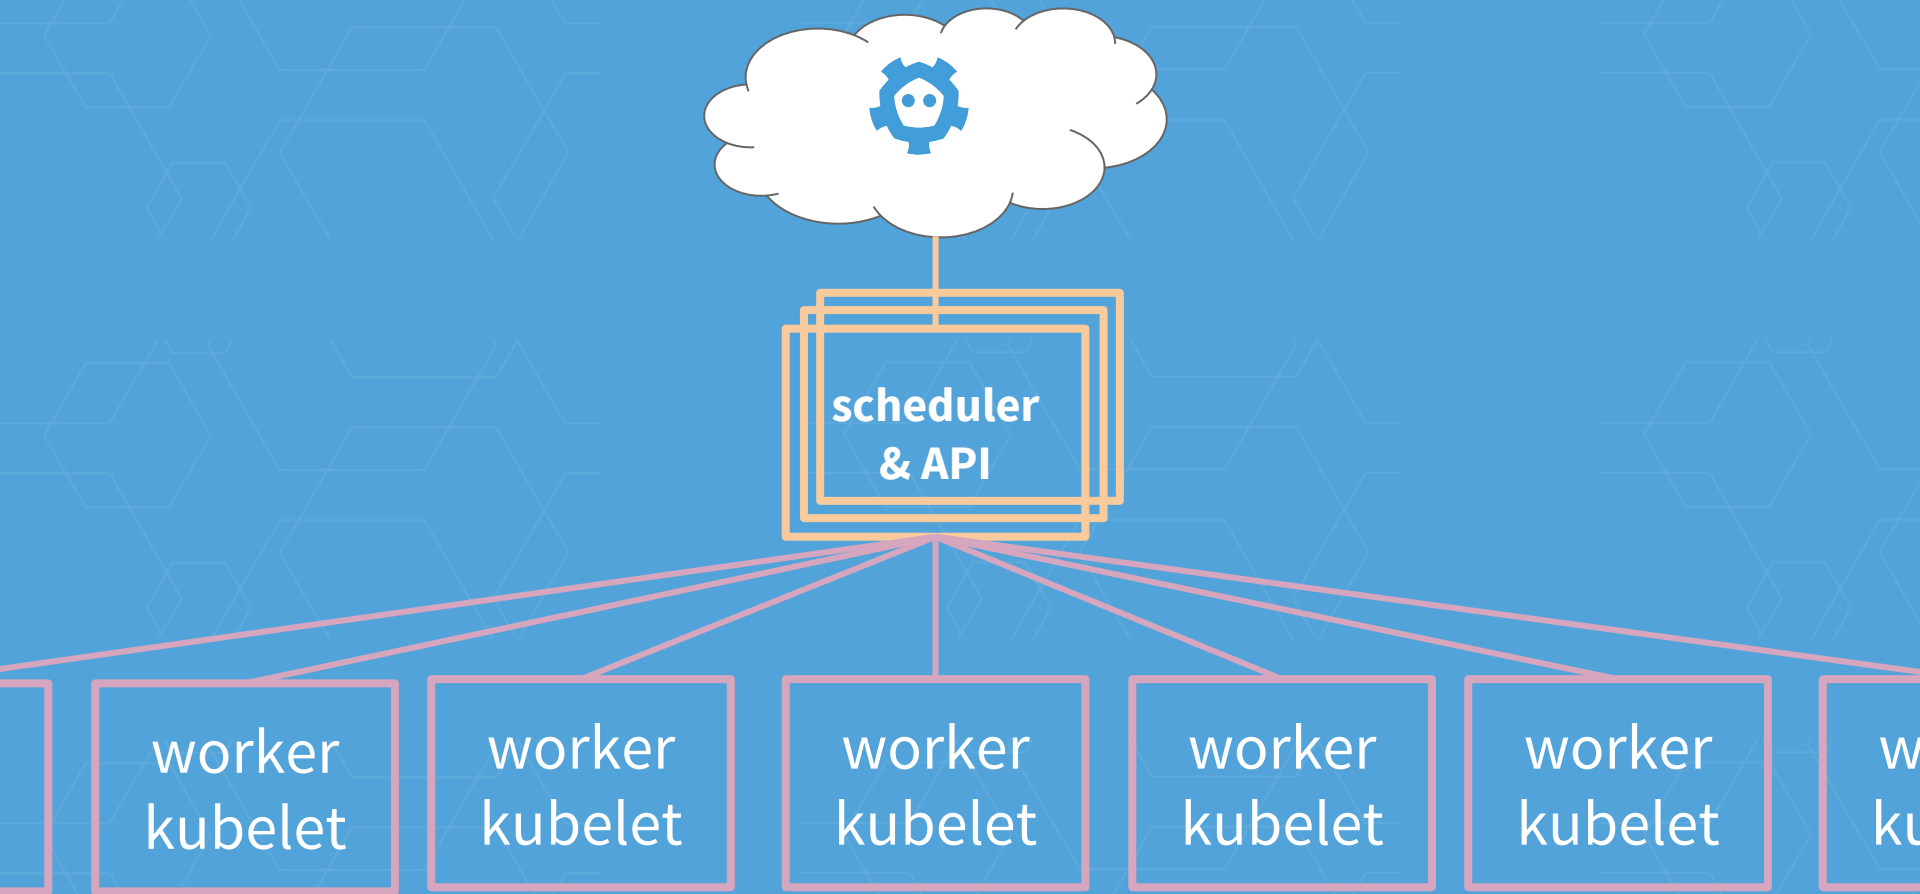
etcd apps
confd

confd

- simple configuration templating
- for "dumb" applications
- watch etcd for changes, render templates with new values, reload applications



etcd apps Kubernetes




Scaling etcd

Scaling etcd to the next level

- Recent improvements
 - Asynchronous snapshots
 - Request pipelining
- Future improvements
 - v3 and beyond


Recent improvements (v2)

- Asynchronous snapshotting
 - append-only log-based system
 - with more and more writes, will grow indefinitely...
 - so, snapshot, purge log




Recent improvements (v2)

- Asynchronous snapshotting
 - before: stop the world, save snapshot to disk
 - problem: this blocks all writes
 - can only resume writing when snapshot is done




Recent improvements (v2)

- Asynchronous snapshotting
 - before: stop the world, save snapshot to disk
 - problem: this blocks all writes
 - can only resume writing when snapshot is done




Recent improvements (v2)

- Asynchronous snapshotting
 - now: in-memory copy, write copy to disk
 - can continue serving writes while snapshotting



Recent improvements (v2)

- Asynchronous snapshotting
 - now: in-memory copy, write copy to disk
 - can continue serving writes while snapshotting



Recent improvements (v2)


- Asynchronous snapshotting
 - now: in-memory copy, write copy to disk
 - can continue serving writes while snapshotting



Recent improvements (v2)


- Asynchronous snapshotting in the real world
 - Discovery service (production etcd cluster serving <https://discovery.etcd.io>)
 - *500 Internal Server Error* spikes dropped to zero

etcd was upgraded




Recent improvements (v2)

- raft pipelining
 - raft is based around a series of RPCs (e.g. AppendEntry)
 - etcd previously used *synchronous RPCs*
 - send next message only after getting previous reply
 - safe, but low throughput



Recent improvements (v2)

- raft pipelining
 - etcd now uses *RPC pipelining*
 - send series of messages without waiting for replies
 - optimistic, recover from failures



Future improvements (v3)

"Scaling etcd to thousands of nodes"

etcd v3.0

- Efficient and powerful API
- Disk-backed storage
- Incremental snapshots

The future (v3)

- Efficient and powerful API
 - flat binary key space
 - multi-object transaction
 - native leasing API
 - native locking API
 - gRPC (HTTP2 + protobuf)

Key space

- Flat binary key-value space
 - coreos=awesome
 - coreos/etcd=kv
 - coreos/rkt=container
- Keep it as simple as possible
 - want hierarchy?
 - build your own layer on top of kv

v3 API

- Put
 - foo=bar
- Get
- Range (consistent multi-get)
 - single key: foo
 - prefix: foo->fop (exclude)
 - range: foo->foo1
- Delete Range
 - same as range

KV API

```
KV.Put("foo", "bar")
```

```
KV.Get("foo")
```

```
KV.Range("foo", "foo10")
```

```
KV.Delete("foo")
```

```
KV.DeleteRange("foo", "foo10")
```

v3 API

- Mini transaction

- two phases
 - *compare*
 - *execution* (either *success* or *failure*)
- compare on value, index, etc.
- execute a list of basic operations

v3 API

- Mini transaction

- compare and swap
 - *compare*: foo=bar
 - *success*: foo=bar2
- multiple object transaction
 - *compare*: cond1=true && cond2=true
 - *success*: pass=true
 - *failure*: pass=false

Mini Transaction

```
Tx.If(  
    Compare(Value("foo"), ">", "bar"),  
    Compare(Version("foo"), "=", 2),  
    ...  
) .Then(  
    OpPut("ok", "true") ...  
) .Else(  
    OpPut("ok", "false") ...  
) .Commit()
```

v3 API

- Watch

- support multiple keys and prefixes per stream
 - `watchKey(foo)`
 - `watchPrefix(coreos)`
- support watch from historical point
 - `watchKey(foo, index_of_an_hour_ago)`
 - user-driven history compaction

gRPC

- Efficient
 - multiple streams share one TCP connection
 - compacted encoding format (protobuf)
- Rich generated libraries in tens of languages
 - Go, Java, Python, C++...

The future (v3)

- Incremental snapshot
 - only save the delta instead of full data set
 - less I/O and CPU cost per snapshot
 - no bursty resource usage, more stable performance

The future (v3)

- Disk backend
 - keep the cold historical data on disk
 - keep the hot data in memory
 - support "entire history" watches
 - user-facing compaction API

The future (v3)

- Upstream recipes for common usage patterns
 - Leases: attaching *ownership* to keys
 - Leader election
 - Locking resources

Lease

```
l := lease.Create(10*second)
```

```
kv.Put("foo", "bar", l.ID)
```

```
// key will be removed without keeping
```

```
// alive the lease
```

```
go KeepAlive(l.id)
```

Watch

```
events, err := watcher.Watch("foo")
if err != nil {
    // handle error
}
for r := <- responses {
    // consume received events
}
```

Recipes

Leader election

- `election.Elect("eFoo"), election.Resign("eFoo")`

Locking

- `locking.Acquire("lFoo"), locking.Release("lFoo")`



coreos.com/fest

May 9 & 10, 2016 | Berlin, Germany

- Early bird tickets
- Sponsorships are still available
- Submit a talk before February 29th!

Thanks!



Join us!

github.com/coreos/etcd

We are hiring!

coreos.com/careers

extra slides
(did I talk too fast?)

MVCC

Rev 0



Revision




MVCC

`KV.Put("foo", "bar")` -> increase Rev

Rev 1

Rev 0



MVCC

KV.Get("foo") = "bar"

Rev 1

Rev 0

foo = bar

Revision



MVCC


```
KV.Get("foo", WithRev(0)) = nil
```



Compaction



Why?

Too many revisions



Compaction

Analyse the old revisions to be compacted



Compaction



Rule 1: the key with tombstone can be removed

Rule 2: keep the latest version of a non-tombstone key




Compaction

Clean up old revisions in background




Compaction

Clean up old revisions in background



Compaction

Done!





etcd basics

fault tolerance

Available



Leader



Follower



Available




Leader




Follower



Available



Unavailable





etcd basics

leader fault tolerance

Available



Leader



Follower



Available




Leader




Follower



Temporarily Unavailable



Available



Unavailable



etcd durability

wal, snapshots, backups

etcd bootstrap

discovery, static

```
$ curl discovery.etcd.io/new?size=5  
https://discovery.etcd.io/6eadeac2...
```

discovery



discovery



discovery




discovery



Leader



Follower



etcd reconfig

live addition and removal



