



the cornerstore of distributed systems
(built using Go)

Jonathan Boule

CoreOS
@baronboule

What is etcd?

/etc
distributed

hence, the name...

a clustered **key-value store**

GET and SET operations

a **building block** for higher order systems

primitives for building reliable
distributed systems

What is etcd?

- distributed /etc
 - cluster-level configuration
- clustered key-value store
- primitives for building reliable distributed systems
 - distributed locking system
 - distributed scheduling system
- supports a lot of large distributed applications
 - SkyDNS, Kubernetes, CloudFoundry, ...

History of etcd

History of etcd

- 2013.8 Alpha release
 - v0.1-v0.4
- 2015.2 Stable release
 - v2.0+
 - stable replication engine (new Raft implementation)
 - stable v2 API
- 2016.?
 - v3.0
 - efficient, powerful API
 - highly scalable backend

History of etcd

- Production-ready!
 - long running failure injection tests
 - no known data loss issues
 - no known inconsistency issues
 - used in critical CoreOS systems like locksmith and fleet
 - trusted by Google, Pivotal, compose and many more!

Why build etcd?

Why build etcd?

- CoreOS: "Secure the internet"
- Updating servers = rebooting servers
- Move towards new application container paradigm
- Need a:
 - shared configuration store (for service discovery etc)
 - distributed lock manager (to co-ordinate reboots)
- Existing solutions were inflexible (undocumented binary API), difficult to configure

Why use etcd?

Why use etcd?

- Highly available
- Highly reliable
- Strong consistency guarantees
- Simple, fast HTTP API
- Open source

"For the most critical data of a distributed system"

**How does etcd
work?**

How does etcd work?

- Replicated log to model a state machine
- Raft: *"In Search of an Understandable Consensus Algorithm"* (Ongaro, 2014)
- Three key concepts
 - Leaders
 - Elections
 - Terms
- etcd clusters elect a leader; all state changes performed by that leader

How does etcd work?

- Go
- `/bin/etcd`
 - daemon
 - 2379 (client requests/HTTP + JSON API)
 - 2380 (peer-to-peer/HTTP + protobuf)
- `/bin/etcdctl`
 - command line client
 - `net/http`, `encoding/json`, ...


etcd basics

clusters


Available




Available



Available



Unavailable



etcd basics

API

Simple HTTP API (v2)

- GET `/v2/keys/foo`
 - Get the value of a key
- GET `/v2/keys/foo?wait=true`
 - Wait for changes on key foo
- PUT `/v2/keys/foo -d value=bar`
 - Set the value of a key
- DELETE `/v2/keys/foo`
 - Delete a key

Compare-and-Swap

```
PUT /v2/keys/foo?prevValue=bar -d  
value=ok
```

```
CAS(/foo, bar, ok)
```

```
if /foo == bar  
    set(/foo, ok)  
else  
    do nothing
```


Compare-and-Delete

DELETE /v2/keys/foo?prevValue=bar

```
CAD(/foo, bar)
```

```
if /foo == bar  
    delete(/foo)  
else  
    do nothing
```

Simple HTTP API (v2)

Native Go bindings

```
import "github.com/coreos/etcd/client"  
  
cl := client.New(client.Config{})  
kapi := client.NewKeysAPI(cl)  
kapi.Set("foo", "bar", ...)
```



etcd apps


etcd apps

locksmith

locksmith

- cluster wide reboot lock
 - "semaphore for reboots"
- CoreOS updates happen *automatically*
 - stop all the machines restarting at once...

Cluster Wide Reboot Lock



needs reboot

Cluster Wide Reboot Lock

- Need to reboot? Decrement the semaphore key (*atomically*) with etcd.
- `manager.Reboot()` and wait...
- After reboot, increment the semaphore key in etcd (*atomically*).

Cluster Wide Reboot Lock

Sem=1




server1


server2

server3

Cluster Wide Reboot Lock





Cluster Wide Reboot Lock



Cluster Wide Reboot Lock

Sem=0




Reboot()

Cluster Wide Reboot Lock

Sem=0




needs reboot




Reboot()


Cluster Wide Reboot Lock




Cluster Wide Reboot Lock




Cluster Wide Reboot Lock



Cluster Wide Reboot Lock



Cluster Wide Reboot Lock



etcd apps

skydns

skydns

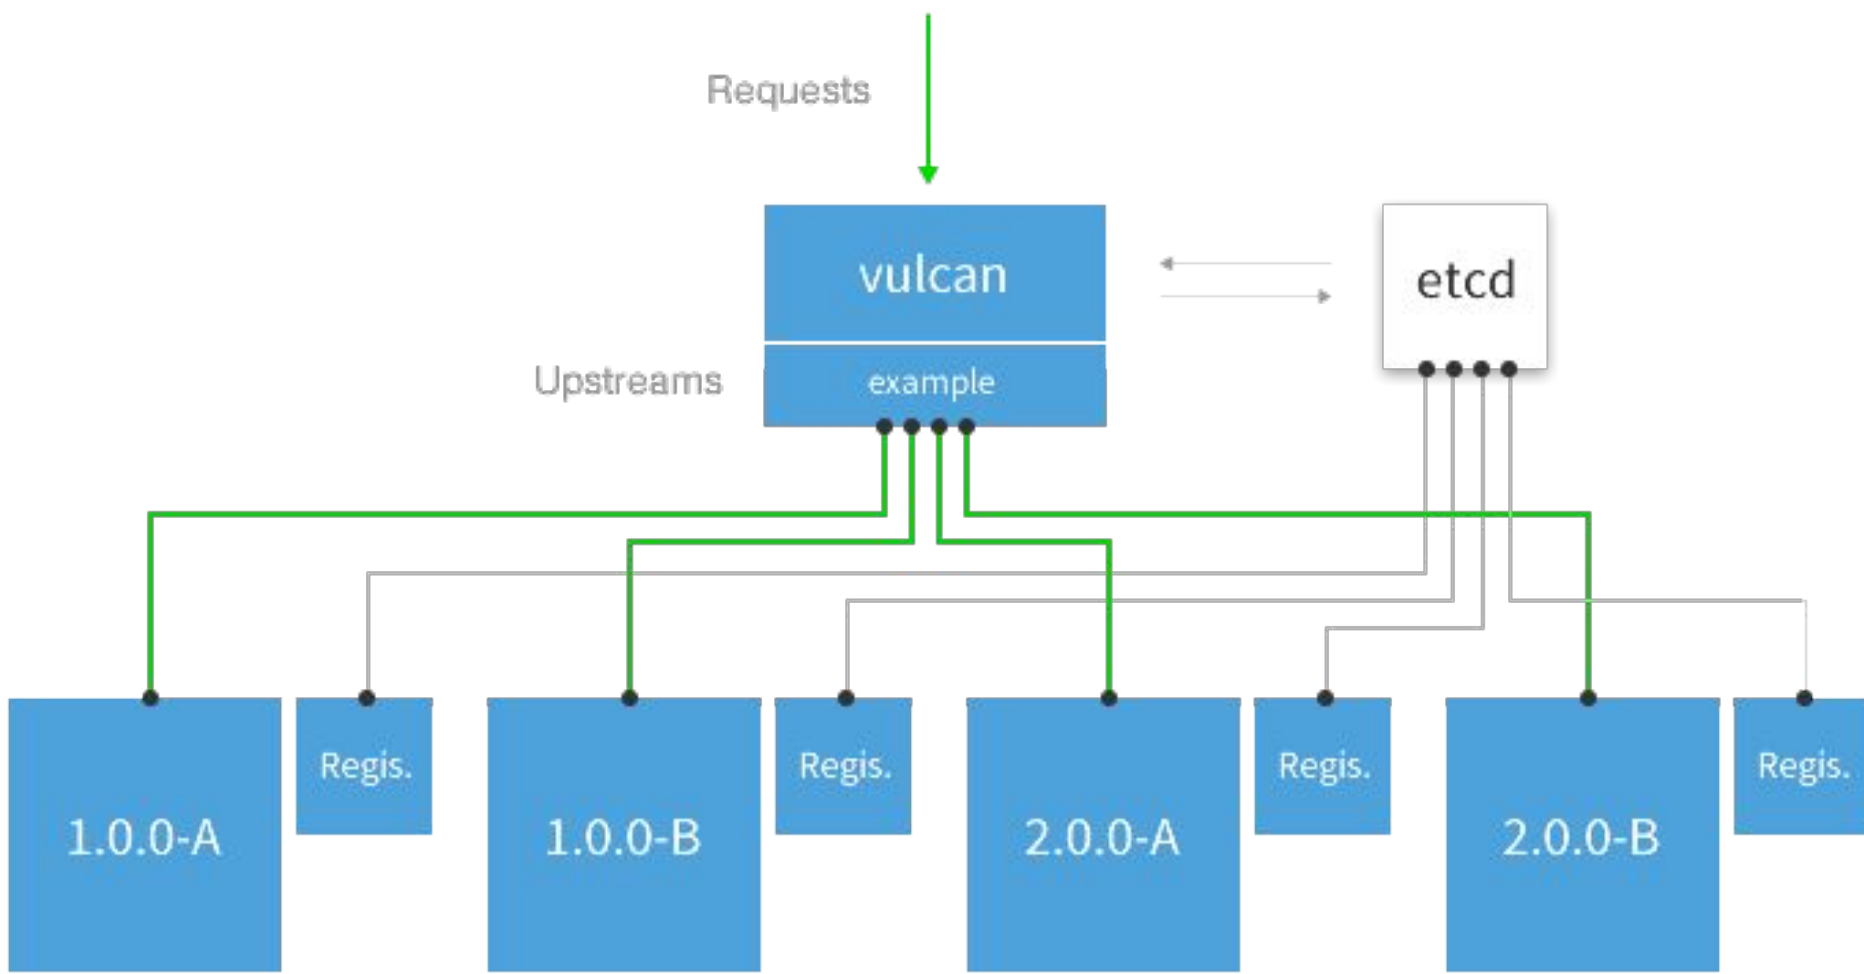
- Service discovery and DNS server
- backed by etcd for all configuration/records

etcd apps

vulcand

vulcand

- "programmatic, extendable proxy for microservices"
- HTTP load balancer
- etcd for all configuration



etcd apps

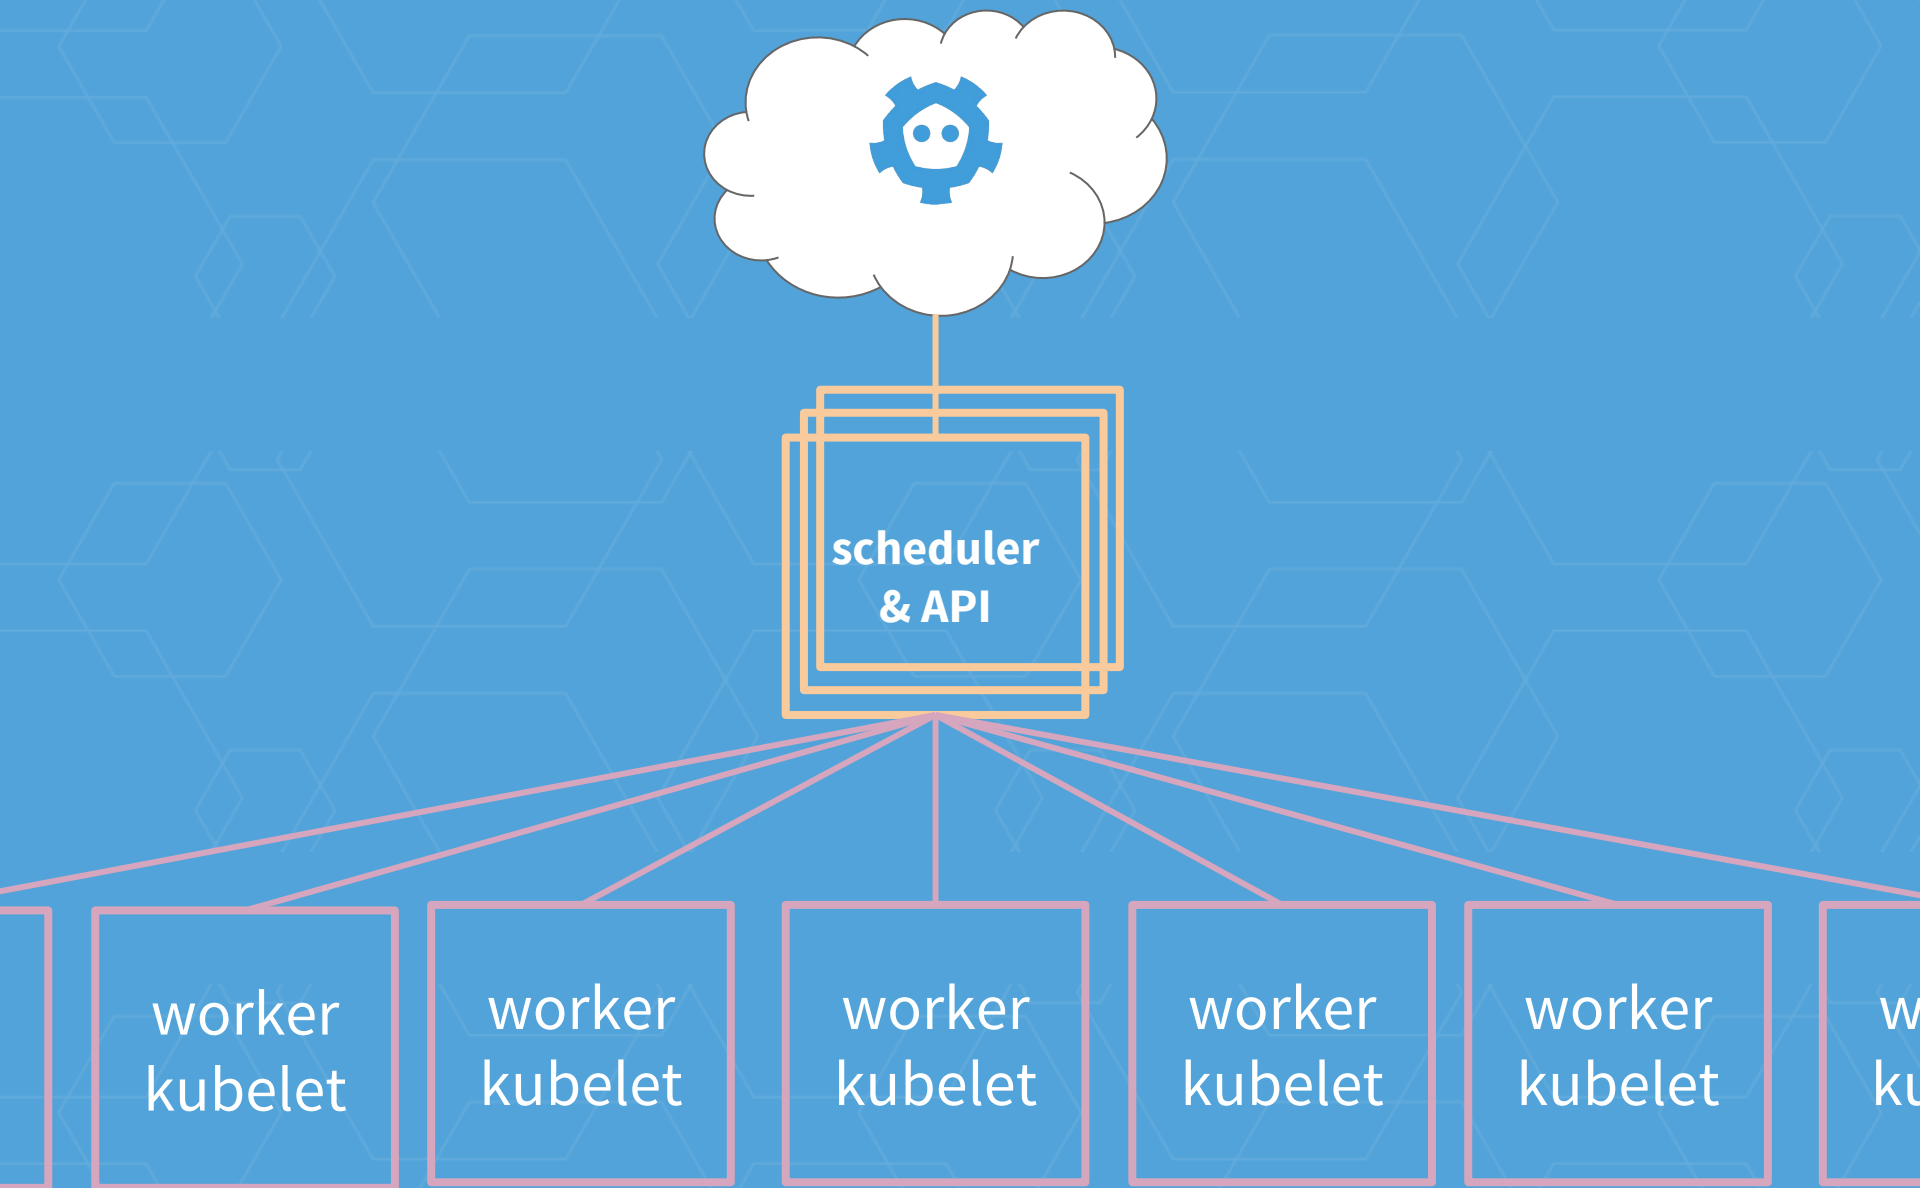
confd

confd

- simple configuration templating
- for "dumb" applications
- watch etcd for changes, render templates with new values, reload applications

etcd apps

Kubernetes




Scaling etcd

Scaling etcd to the next level

- Recent improvements
 - Asynchronous snapshots
 - Request pipelining
- Future improvements
 - v3 and beyond


Recent improvements (v2)

- Asynchronous snapshotting
 - log-based system
 - snapshot before purging log




Recent improvements (v2)

- Asynchronous snapshotting
 - log-based system
 - snapshot before purging log



Recent improvements (v2)

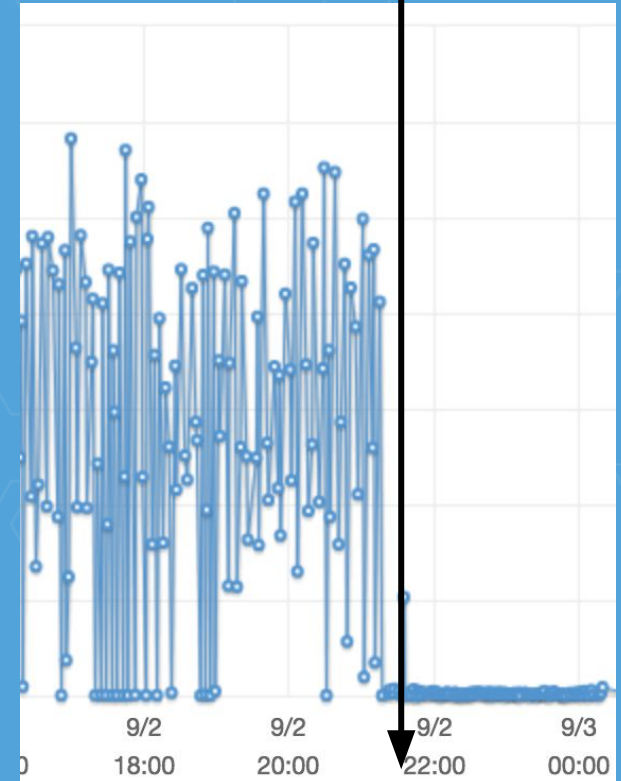
- Asynchronous snapshotting
 - log-based system
 - snapshot before purging log



Recent improvements (v2)


- Asynchronous snapshotting
 - result from discovery service

Upgrade etcd




Recent improvements (v2)

- raft pipelining
 - etcd previously used *synchronous RPCs*
 - send next message only after getting previous reply



Recent improvements (v2)

- raft pipelining
 - etcd now uses *RPC pipelining*
 - send series of messages without waiting for replies



Future improvements (v3)

"Scaling etcd to thousands of nodes"

etcd v3.0

- Efficient and powerful API
- Disk-backed storage
- Incremental snapshots

The future (v3)

- Efficient and powerful API
 - flat binary key space
 - multi-object transaction
 - native leasing API
 - native locking API
 - gRPC (HTTP2 + protobuf)

Key space

- Flat binary key-value space
 - coreos=awesome
 - coreos/etcd=kv
 - coreos/rkt=container
- Keep it as simple as possible
 - want hierarchy?
 - build your own layer on top of kv

v3 API

- Put
 - foo=bar
- Get
- Range (consistent multi-get)
 - single key: foo
 - prefix: foo->fop (exclude)
 - range: foo->foo1
- Delete Range
 - same as range

KV API

```
KV.Put("foo", "bar")
```

```
KV.Get("foo")
```

```
KV.Range("foo", "foo10")
```

```
KV.Delete("foo")
```

```
KV.DeleteRange("foo", "foo10")
```

v3 API

- Mini transaction
 - two phases
 - *compare*
 - *execution* (either *success* or *failure*)
 - compare on value, index, etc.
 - execute a list of basic operations

v3 API

- Mini transaction
 - compare and swap
 - *compare*: foo=bar
 - *success*: foo=bar2
 - multiple object transaction
 - *compare*: cond1=true && cond2=true
 - *success*: pass=true
 - *failure*: pass=false

Mini Transaction

```
Tx.If(  
    Compare(Value("foo"), ">", "bar"),  
    Compare(Version("foo"), "=", 2),  
    ...  
) .Then(  
    OpPut("ok", "true")...  
) .Else(  
    OpPut("ok", "false")...  
) .Commit()
```

v3 API

- Watch
 - support multiple keys and prefixes per stream
 - `watchKey(foo)`
 - `watchPrefix(coreos)`
 - support watch from historical point
 - `watchKey(foo, index_of_an_hour_ago)`
 - user-driven history compaction

gRPC

- Efficient
 - multiple streams share one TCP connection
 - compacted encoding format (protobuf)
- Rich generated libraries in tens of languages
 - Go, Java, Python, C++...

The future (v3)


- Incremental snapshot
 - only save the delta instead of full data set
 - less I/O and CPU cost per snapshot
 - no bursty resource usage, more stable performance

The future (v3)

- Disk backend
 - keep the cold historical data on disk
 - keep the hot data in memory
 - support "entire history" watches
 - user-facing compaction API

MVCC

Rev 0



Revision

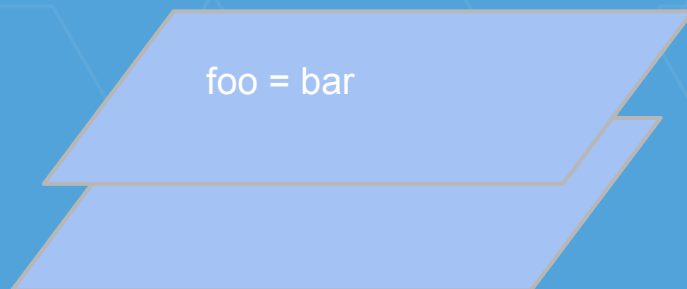


MVCC

`KV.Put("foo", "bar")` -> increase Rev

Rev 1

Rev 0




Revision




MVCC

`KV.Get("foo") = "bar"`



MVCC


```
KV.Get("foo", WithRev(0)) = nil
```



Compaction



Why?

Too many revisions



Compaction


Analyse the old revisions to be compacted



Compaction



Rule 1: the key with tombstone can be removed

Rule 2: keep the latest version of a non-tombstone key





Compaction

Clean up old revisions in background




Compaction

Clean up old revisions in background




Compaction

Done!



KV Snapshot

```
// working on snapshot of KV at Rev 100  
KV.Get("foo", WithRev(100))  
KV.Range("foo", "foo10", WithRev(100))
```



Lease

```
l := lease.Create(10*second)
```

```
kv.Put("foo", "bar", l.ID)
```

```
// key will be removed without keeping  
// alive the lease  
go KeepAlive(l.id)
```

Watch

```
events, err := watcher.Watch("foo")
if err != nil {
    // handle error
}
for r := <- responses {
    // consume received events
}
```

Recipes

Leader election

- `election.Elect("eFoo"), election.Resign("eFoo")`

Locking

- `locking.Acquire("lFoo"), locking.Release("lFoo")`

Barrier

- `barrier.Enter("bFoo"), barrier.Leave("bFoo")`

etcd and go

etcd and go: the good

- Extremely fast development speed
- Generally robust standard libraries
- Healthy, active ecosystem
- Simple but powerful concurrency

... i.e., all the usual reasons people like Go

etcd and go: the less good

- HTTP pipelining + CloseNotify = no fun
 - <http://www.projectclearwater.org/adventures-in-debugging-etcd-http-pipelining-and-file-descriptor-leaks/>
 - <https://github.com/golang/go/issues/13165>
- CloseNotify = no fun
 - <https://github.com/golang/go/issues/9524>
- Unpredictable GC = latency spikes
 - <https://github.com/coreos/etcd/issues/4111>
- Scheduler starvation = dead Raft goroutine
 - citation needed

Thanks!



Join us!

github.com/coreos/etcd



We are hiring!

coreos.com/careers