

# rkt and Kubernetes

What's new (and coming) with  
Container Runtimes and Orchestration

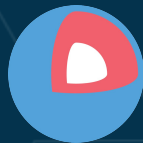




# Jonathan Boulle

[github.com/jonboulle](https://github.com/jonboulle) - [@baronboulle](https://twitter.com/baronboulle)

---



Core OS

# Why rkt and Kubernetes?

~~Why rkt and Kubernetes?~~  
**Why container runtimes  
and orchestration?**

# CoreOS, Inc (2013 - today)

Mission: "Secure the Internet"

## Started at the OS level: CoreOS Linux

- Modern, minimal operating system
- Self-updating (read-only) image
- Updates must be *automatic* and *seamless*

# *Automatic and seamless*

- If the OS is always updating, what about applications running on it?
- Classic use case for *containers* and *orchestration*
  - *containers* decouple the application and OS update lifecycles (update at different cadences)
  - *orchestration* decouples application and OS uptime (services can remain unaffected during OS downtime)

# Why container runtimes?

- Update the OS without affecting application dependencies

kernel  
systemd  
rkt  
ssh  
docker

python  
java  
nginx  
mysql  
openssl

stro distro distro distro distro

app



kernel  
systemd  
rkt  
ssh  
docker

stro distro distro distro distro

python  
java  
nginx  
mysql  
openssl

app

kernel  
systemd  
rkt  
ssh  
docker

stro distro distro distro distro

python  
openssl-A

app1

java  
openssl-B

app2

java  
openssl-B

app3

# CoreOS

stro distro distro distro distro


container

container

container

# Why orchestration?

- Update the OS without affecting application uptime



The diagram shows three servers, each represented by a rounded rectangle with a white border. The first server on the left contains three applications (app1, app2, app3) and is labeled 'server1'. The second server in the middle contains two applications (app4, app5) and is labeled 'server2'. The third server on the right contains two applications (app6, app7) and is labeled 'server3'. The background is blue with a faint hexagonal pattern.

app1  
app2  
app3




server1

app4  
app5




server2

app6  
app7

server3




updating...




**needs reboot**

# Without orchestration




**rebooting...**








# Without orchestration



rebooting...




# With orchestration




**needs reboot**

# With orchestration



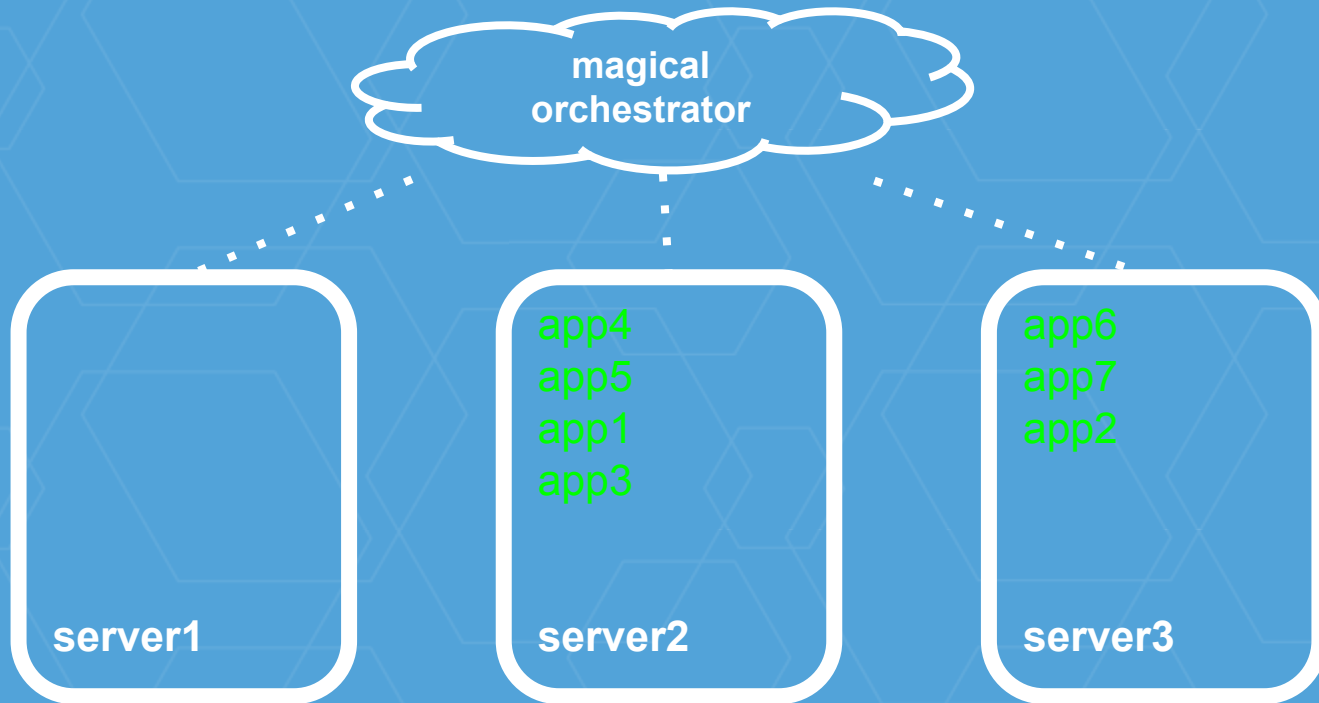
**needs reboot**

# With orchestration




**rebooting...**

# With orchestration




updated!

# With orchestration



updated!

# With orchestration



# Why container runtimes and orchestration?

So we can provide *seamless updates* and push forward the security of application servers



The background is a solid blue color with a repeating pattern of white-outlined hexagons of various sizes, creating a geometric, honeycomb-like texture.

# Why rkt?

**A long time ago in an  
ecosystem far, far away....**

(2014, to be precise)

# Before the "Container Wars"

(please stop saying this)

# 2014

- Large incumbent container tool (in CoreOS)
- *Common* practices, but few *best* practices
  - unsigned images (`curl | sudo sh -`)
  - inefficient/insecure images (`FROM ubuntu:14.04`)
  - PID1 or not to PID1 ([zombie reaping problem](#))
- New platforms emerging, difficult to integrate
  - `systemd + dockerd` = sad times had by all

# 2014 (December)

- Enter rkt (and appc)
  - Create an *alternative* container runtime (competition drives innovation)
  - Emphasise the importance of *security* and *composability*
  - Spur conversation around *standards* in the application container ecosystem



a modern, secure container runtime  
a simple, composable tool  
an implementation of an open standard



a standard application container  
open specification  
associated tooling



# appc

[github.com/appc/spec](https://github.com/appc/spec)

[github.com/appc/acbuild](https://github.com/appc/acbuild)

[github.](https://github.com)

[com/appc/docker2aci](https://github.com/appc/docker2aci)

[github.com/appc/cni](https://github.com/appc/cni)

[github.com/appc/...](https://github.com/appc/...)





# appc

**[github.com/appc/spec](https://github.com/appc/spec) ("appc spec")**

[github.com/appc/acbuild](https://github.com/appc/acbuild)

[github.com/appc/docker2aci](https://github.com/appc/docker2aci)

**[github.com/appc/cni](https://github.com/appc/cni) (more on this later..)**

[github.com/appc/...](https://github.com/appc/)

# appc spec in a nutshell

- Image Format (ACI)
  - what does an application consist of?
- Image Discovery
  - how can an image be located?
- Pods
  - how can applications be grouped and run?
- Executor (runtime)
  - what does the execution environment look like?

# appc spec in a nutshell

- Image Format (ACI)
  - what does an application consist of?
- Image Discovery
  - how can an image be located?
- **Pods**
  - **how can applications be grouped and run?**
- Executor (runtime)
  - what does the execution environment look like?

# appc pods

- grouping of applications executing in a shared context (network, namespaces, volumes)
- shared fate
- the *only* execution primitive: single applications are modelled as singleton pods

# appc pods $\approx$ Kubernetes pods

- grouping of applications executing in a shared context (network, namespaces, volumes)
- shared fate
- the *only* execution primitive: single applications are modelled as singleton pods



a modern, secure container runtime  
a simple, composable tool (CLI)  
an implementation of an open standard (appc)

# 2016

- Docker and rkt both "production-ready"
- Kubernetes too!
- Standards?
  - appc (December 2014)
  - OCI (June 2015)
  - CNCF (December 2015)



appc



OPEN CONTAINER  
INITIATIVE



CLOUD NATIVE  
COMPUTING FOUNDATION

# **rkt architecture**

A quick introduction



# rkt - simple CLI tool

no central daemon

no (mandatory) API

apps run directly under spawning process

bash/systemd/kubelet



rkt run ...



application(s)

# rkt internals

modular architecture  
execution divided into *stages*  
stage0 → stage1 → stage2

# rkt internals

modular architecture  
take advantage of different technologies  
provide a consistent experience to users

bash/systemd/kubelet



rkt run ...



application(s)

bash/systemd/kubelet



rkt run ...



pod


bash/systemd/kubelet... (invoking process)



rkt (stage0)



pod (stage1)



app1 (stage2)

app2 (stage2)

# stage0 (rkt binary)

- primary interface to rkt
- discover, fetch, manage application images
- set up pod filesystems
- manage pod lifecycle
  - rkt run
  - rkt image list
  - rkt gc
  - ...



# stage1 (swappable execution engines)

- default implementation
  - based on systemd-nspawn+systemd
  - Linux namespaces + cgroups for isolation
- kvm implementation
  - based on lkvm+systemd
  - hardware virtualisation for isolation
- others?
  - e.g. xhyve (OS X), unc (unprivileged containers)

## stage2 (inside the pod)

- actual app execution
- independent filesystems (chroot)
- shared namespaces, volumes, IPC, ...

# rkt TPM measurement (new!)

- TPM, Trusted Platform Module
  - physical chip on the motherboard
  - cryptographic keys + processor
- Used to "*measure*" system state
- Historically just use to verify bootloader/OS (on proprietary systems)

# rkt TPM measurement (new!)

- CoreOS added support to GNU Grub
- rkt can now record information about running pods in the TPM
- attestable record of what images and pods are running on a system

# rkt TPM measurement (new!)

Containers

**rkt**

Verify images with trusted keys

Verify configuration state

OS

**CoreOS Linux**

Verify integrity of the OS release

Hardware

**Firmware & TPM**

Customer key embedded in firmware

Tamper-proof  
Audit log

# rkt TPM measurement (new!)

- For much, much more on TPM and rkt, see Matthew Garrett's talk:  
*"Integrated trusted computing in Kubernetes"*
- 11:30am tomorrow

# rkt API service (new!)

- *optional*, gRPC-based API daemon
- exposes *read-only* information on pods/images
- runs as *unprivileged* user
- easier integration with other projects

# Why rkt?

Secure  
Standards  
Composable





+



kubernetes

# rkt + Kubernetes


rkt ♥ k8s in a few ways:

- using rkt as container runtime (aka "*rktnetes*")
- using rkt to run Kubernetes ("*rkt fly*")
- integrating with rkt networking (CNI)

# Kubelet + Container Runtimes

- Kubelet provides a Runtime interface
  - SyncPod()
  - GetPod()
  - KillPod()
  - ...
- in theory, anyone can implement this
- in practise, lots of Docker assumptions

# Kubelet + Docker (default)



# Kubelet + Docker (default)

## Problems:


- Docker doesn't understand *Pods*
  - kubelet must maintain pod<->container mapping
  - "infra container" to hold namespaces for pod
- dockerd = SPOF for node
  - if Docker goes down, so do all containers
- Docker doesn't interact well with systemd

# Kubelet + rkt (rktnetes)


Using rkt as the kubelet's container runtime

- *A pod-native* runtime
- First-class integration with systemd hosts
- *self-contained pods* process model = no SPOF
- Multi-image compatibility (e.g. docker2aci)
- Transparently swappable

# Kubelet + rkt (rktnetes - with systemd)



# Kubelet + rkt (rktnetes - without systemd)





# rktnetes today

Nearly complete!  
80% of end-to-end tests passing  
cAdvisor integration in progress

# Using rkt to run Kubernetes

- Kubernetes components are largely self-hosting, but not entirely
  - Need a way to bootstrap kubelet on the host
  - kubelets can then host control plane components
- On CoreOS, this means in a container..
  - ... but kubelet has some unique requirements (like mounting volumes on the host)

# Using rkt to run Kubernetes

- rkt "*fly*" feature (new in 0.15.0+)
- unlike rkt run, does *\*not\** execute pods
- execute a *single application* in an *unconstrained environment*
- all the other advantages of rkt (image discovery, signing/verification, management)

bash/systemd/... (invoking process)

└─▶ rkt (stage0) - without *fly*

└─▶ pod (stage1)

└─▶ app1 (stage2)

└─▶ app2 (stage2)

bash/systemd/... (invoking process)

└─▶ rkt (stage0) - without *fly*

└─▶ pod (stage1)

└─▶ app1 (stage2)

└─▶ app2 (stage2)

*Isolated mount (and PID, ...) namespace*

bash/systemd/... (invoking process)

└─▶ rkt (stage0) - with *fly*

└─▶ application

bash/systemd/... (invoking process)

└─▶ rkt (stage0) - with *fly*

└─▶ application

*Host mount (and PID, ...) namespace*

bash/systemd/... (invoking process)

└─▶ rkt (stage0) - with *fly*

└─▶ kubelet

*Host mount (and PID, ...) namespace*



# rkt networking

Plugin-based

IP(s)-per-pod

Container Networking Interface (CNI)

# Container Runtime (e.g. rkt)

---

## Container Networking Interface (CNI)

---

**ptp**

**macvlan**

**ipvlan**

**OVS**

# CNI in a nutshell


- Container can join multiple networks
- Network described by JSON config
- Plugin supports two commands
  - ADD container to the network
  - REMOVE container from the network
- Plugins are responsible for all logic
  - allocating IPs, talking to backend components, ...

# CNI: example configuration

```
{  
  "name": "mynet",  
  "type": "ptp",  
  "ipam": {  
    "type": "host-local",  
    "subnet": "10.1.1.0/24"  
  }  
}
```

```
$ rkt run --net=mynet coreos.com/etcd
```

# How rkt uses CNI



`/var/lib/rkt/pods/run/$POD_UUID/netns`

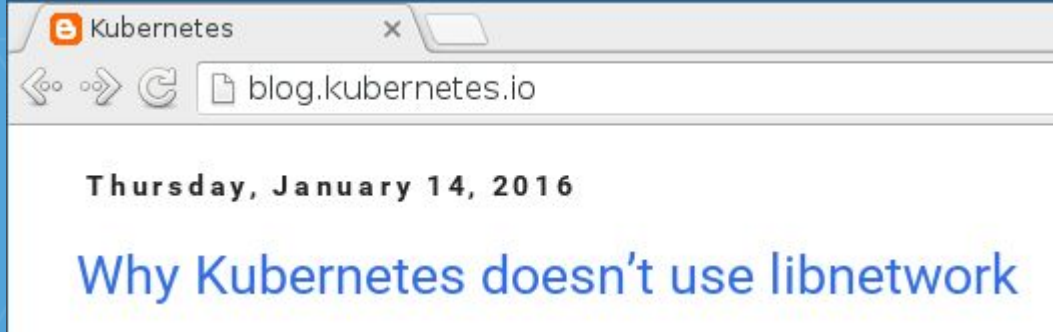
# Kubernetes networking

Plugin-based (but never left alpha)

IP(s)-per-pod  
(sound familiar?)

# Kubernetes and CNI

Previously CNI was just another plugin type, but soon to be "*the* Kubernetes plugin model"



# CNI today

v0.x

Handles all networking in rkt  
Integrations with Project Calico, Weaveworks



# Looking ahead

What's coming up for rkt and Kubernetes

# rktnetes 1.0

end of 2016Q1

Fully supported, full feature parity  
Automated end-to-end testing on CoreOS

# rktnetes 1.0+

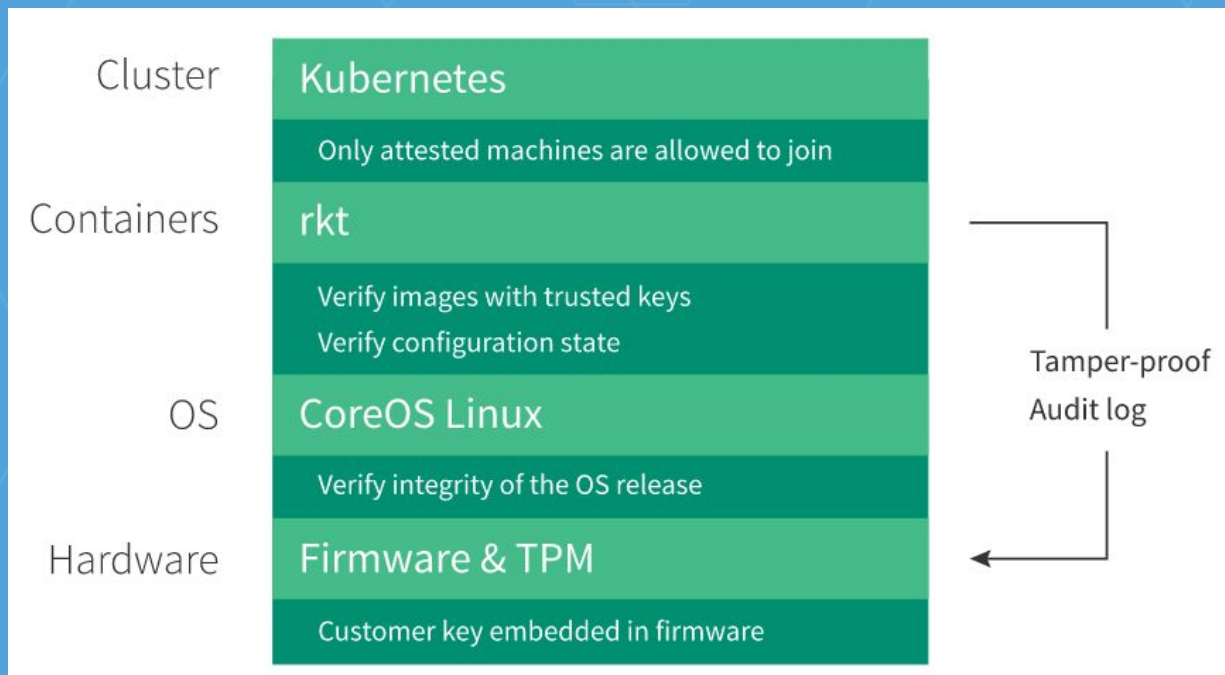
LKVM backend by default

Native support for ACI\* in Kubernetes API

TPM up to the Kubernetes level

(\*and/or OCI, ...)

# Tectonic Trusted Computing



<https://coreos.com/blog/coreos-trusted-computing.html>

# CNI 1.0

Stable configuration

Stateless plugins (runtime responsibility)

IPv6

*<your suggestions here>*

# Kubernetes 1.3

Move all networking code into CNI plugins

# Kubelet upgrades

- Remember from CoreOS mission:  
"updates must be *automatic and seamless*"
- If kubelet is in OS, must be upgraded in lock-step
- But mixed-version clusters don't always work  
(e.g. upgrading from 1.07 - 1.1.1: <https://github.com/kubernetes/kubernetes/issues/16961> )

# Kubelet upgrades

- Solution: API driven upgrades
- Small agent living on host, invoking kubelet (using rkt fly)
- Reading annotations from the kubelet API server
- Follow along:

<https://github.com/coreos/bugs/issues/1051>



# Graceful kubelet shutdown

- When an update is ready, locksmith signals kubelet to gracefully shut down
- Kubernetes can then gracefully migrate apps before shutdown
- <https://github.com/coreos/bugs/issues/1112>
- <https://github.com/kubernetes/kubernetes/issues/7351>

# tl;dr:

- Use rkt
- Use Kubernetes
- Use rkt + Kubernetes (rktnetes)
- Get involved and help define the future of application containers and Kubernetes



**May 9 & 10, 2016 - Berlin, Germany**

**[coreos.com/fest](http://coreos.com/fest) - [@coreosfest](https://twitter.com/coreosfest)**

# Questions?



kubernetes



CoreOS

## Join us!

contribute: [github.com/coreos/rkt](https://github.com/coreos/rkt)

careers: [coreos.com/careers](https://coreos.com/careers) *(now in Berlin!)*

# rkt security ("secure by default")

- image signature verification
- privilege separation
  - e.g. fetch images, expose API (new!) as non-root
- SELinux integration (although </3 overlayfs..)
- lkvm stage1 for true hardware isolation
- TPM attestation (new!)