


# Metrics Matter

Debugging the Kubernetes scheduler with Prometheus  
(aka, 10x improvement for Kubernetes performance)

Jonathan Boule, CoreOS

 [github.com/jonboulle](https://github.com/jonboulle)

 [@baronboulle](https://twitter.com/baronboulle)



# **Why metrics?**

Why "instrument all the things"?

# **Health monitoring** with metrics

Factors externally visible to your system: request/error rates, latencies

# **Alerting** with metrics

Time-series based alerting, same language and power as monitoring

**What else can we do?**

# **Business insight** with metrics

Number of page views, capacity planning, ...

# **Autotuning** with metrics

Automatically adjust system state based on metrics

# **Debugging** with metrics

Instrument everything, expose deep application internals



# **Debugging: a case study**

Improving the Kubernetes scheduler performance by 10x

# **Debugging: a case study**

How to use Prometheus

# **Debugging: a case study**

How to ~~not~~ use Prometheus

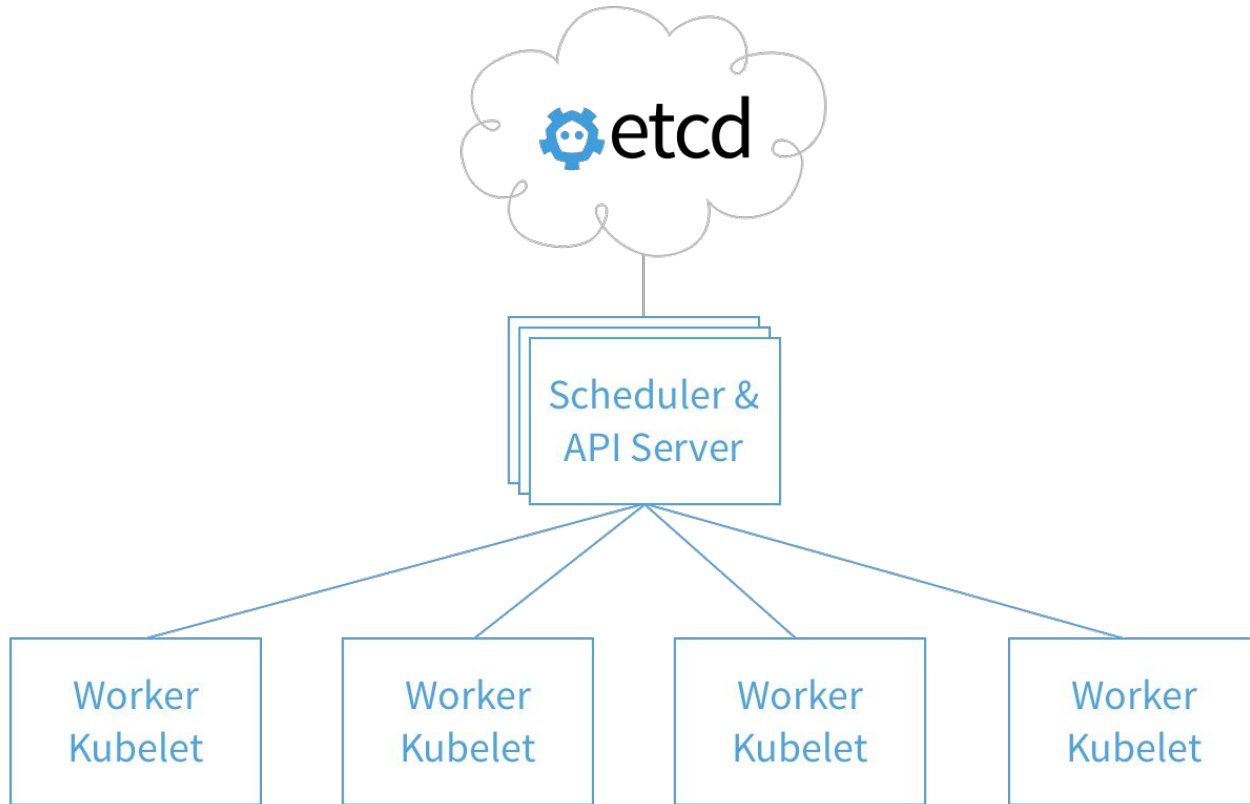
**Initial goal: understand Kubernetes performance**

New "Kubernetes scalability" team at CoreOS

Test and understand Kubernetes cluster performance

(ultimate goal: ensure it can scale to 10 000s of nodes)

# Kubernetes architecture overview



# Initial focus: API server and scheduler (control plane)

Worker nodes are relatively independent

Control plane is involved in every single pod scheduled in a cluster

# Initial experiments: density tests

**Kubemark:** simulate the worker nodes, test large clusters

How long does it take to run 30 pods on every node in a cluster?

Start with a 100 node cluster => 3000 pods to schedule

# Initial experiments: density tests

**Kubemark:** simulate the worker nodes, test large clusters

How long does it take to run 30 pods on every node in a cluster?

Start with a 100 node cluster => 3000 pods to schedule

```
Pods: 229 out of 3000 created, 211 running, 18 pending, 0 waiting
Pods: 429 out of 3000 created, 412 running, 17 pending, 0 waiting
Pods: 622 out of 3000 created, 604 running, 17 pending, 1 waiting
...
Pods: 2578 out of 3000 created, 2561 running, 17 pending, 0 waiting
Pods: 2779 out of 3000 created, 2761 running, 18 pending, 0 waiting
Pods: 2979 out of 3000 created, 2962 running, 16 pending, 1 waiting
Pods: 3000 out of 3000 created, 3000 running, 0 pending, 0 waiting
```




# Initial results

3000 pods in ~150 seconds

Linear rate of pod **creation**  
(20 pods/second)

No delay between creation+running

Suspected bottleneck in overall  
performance; starting point for  
further investigation



# Initial suspect: hardware resource starvation

First place to look

Easy fix: throw more hardware at the problem!

# Initial suspect: hardware resource starvation

Using 3 cores out of 16 cores

PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
20	0	268m	253m	23m	S	192	0.4	4:25.43	kube-apiserver
20	0	58000	47m	15m	S	53	0.1	0:51.18	kube-scheduler
20	0	130m	118m	20m	S	17	0.2	0:27.61	kube-controller
20	0	124m	111m	10m	S	15	0.2	0:21.31	etcd

*Resource usage during the test*

# ~~Initial suspect: hardware resource starvation~~

Using 3 cores out of 16 cores

PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
20	0	268m	253m	23m	S	192	0.4	4:25.43	kube-apiserver
20	0	58000	47m	15m	S	53	0.1	0:51.18	kube-scheduler
20	0	130m	118m	20m	S	17	0.2	0:27.61	kube-controller
20	0	124m	111m	10m	S	15	0.2	0:21.31	etcd

*Resource usage during the test*

~~Initial suspect: hardware resource starvation~~

Using 3 cores out of 16 cores

PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
20	0	268m	253m	23m	S	192	0.4	4:25.43	kube-apiserver
20	0	58000	47m	15m	S	53	0.1	0:51.18	kube-scheduler
20	0	130m	118m	20m	S	17	0.2	0:27.61	kube-controller
20	0	124m	111m	10m	S	15	0.2	0:21.31	etcd

*Resource usage during the test*

We'll come back to this...

~~Initial suspect: hardware resource starvation~~

Bottleneck somewhere in the software

But... finding bottlenecks in a complex codebase is never trivial

(and Kubernetes is large: ~1.3M LOC)

# Prometheus to the rescue!

Kubernetes provides metrics for most end-to-end API calls

Use Prometheus to explore suspicious areas - narrow down the focus of the investigation

# Narrowing it down

Re-ran the Kubemark suite, monitoring scheduler's metrics

```
# HELP scheduler_e2e_scheduling_latency_microseconds E2e scheduling latency (scheduling algorithm + binding)
# TYPE scheduler_e2e_scheduling_latency_microseconds summary
scheduler_e2e_scheduling_latency_microseconds{quantile="0.5"} 23207
scheduler_e2e_scheduling_latency_microseconds{quantile="0.9"} 35112
scheduler_e2e_scheduling_latency_microseconds{quantile="0.99"} 40268
scheduler_e2e_scheduling_latency_microseconds_sum 7.1321295e+07
```



# Narrowing it down

Re-ran the Kubemark suite, monitoring scheduler's metrics

```
# HELP scheduler_e2e_scheduling_latency_microseconds E2e scheduling latency (scheduling algorithm + binding)
# TYPE scheduler_e2e_scheduling_latency_microseconds summary
scheduler_e2e_scheduling_latency_microseconds{quantile="0.5"} 23207
scheduler_e2e_scheduling_latency_microseconds{quantile="0.9"} 35112
scheduler_e2e_scheduling_latency_microseconds{quantile="0.99"} 40268
scheduler_e2e_scheduling_latency_microseconds_sum 7.1321295e+07
```

Scheduler end-to-end latency ~7ms => ~140 pods/second

Implied that scheduler was not the (primary) bottleneck

# Narrowing it down

Explore more metrics, logs (not enough exposed metrics!)

Thanks to log debugging, isolated contention to Replication  
Controller code

# Iterative improvements

Removed rate limiters (protecting the API from being overwhelmed)

Fixed an inefficient code path....

# Inefficient code path

	51	+func init() {
	52	+       metrics.Register()
	53	+}
	54	+
51	55	// HTTPClient is an interface for testing a request object
52	56	type HTTPClient interface {
53	57	Do(req *http.Request) (*http.Response, error)
		@@ -109,7 +113,6 @@ type Request struct {
109	113	// NewRequest creates a new request helper object for accessing
110	114	func NewRequest(client HTTPClient, verb string, baseURL *url.URL,
111	115	codec runtime.Codec) *Request {
112		-       metrics.Register()


# Iterative improvements

Removed rate limiters (protecting the API from being overwhelmed)

Fixed an inefficient code path....

Early success!

From **20 pods/sec** to **300+ pods/second** (100-node cluster)



*Increased pod scheduling throughput from our changes*

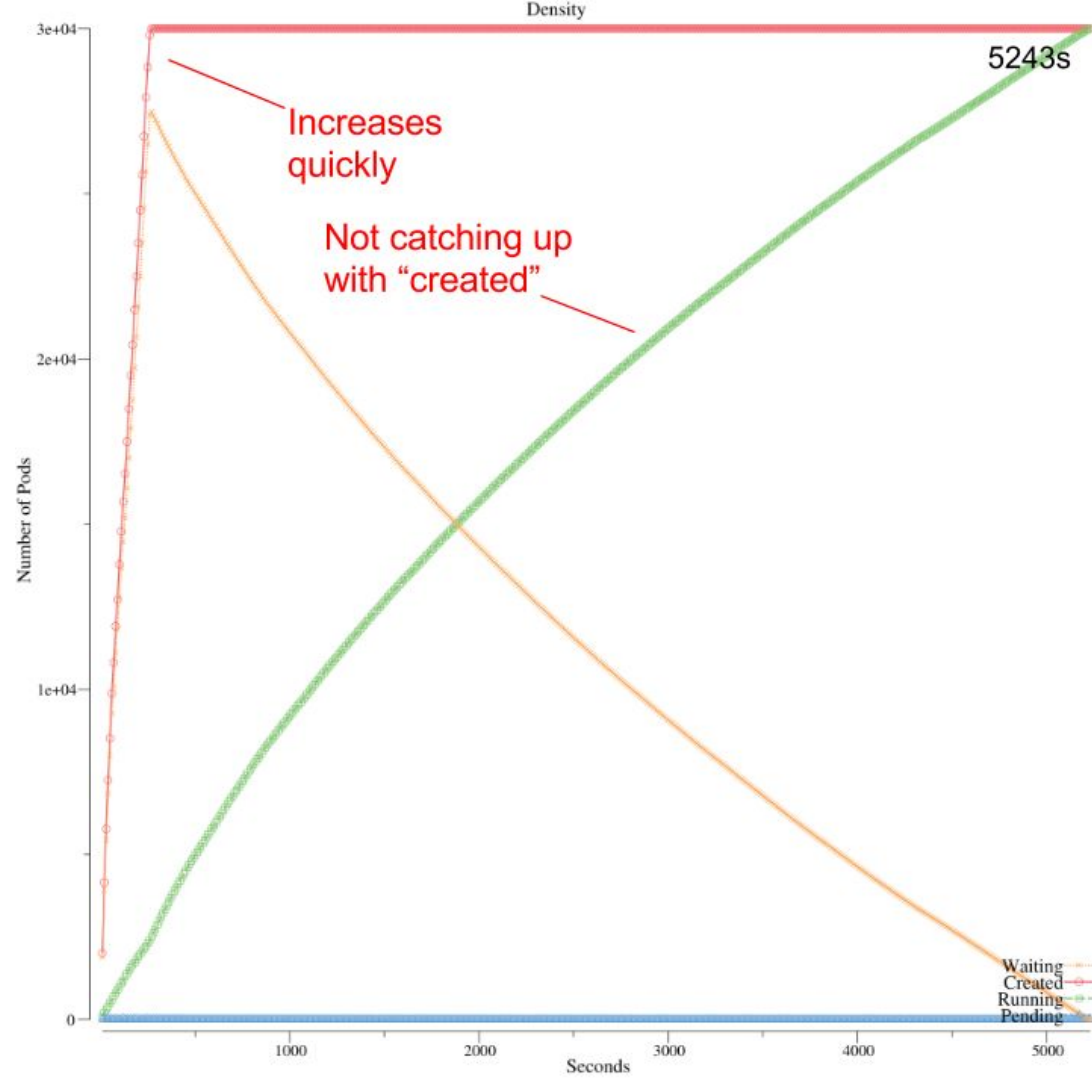
# Expanding horizons

Retried tests on a 1 000-node cluster

Not so lucky: **5 243 seconds** to run **30 000 pods**

Average pod throughput of **5.72 pods/sec**

Creation rate was good, but running rate remained low





# Changing focus: the scheduler

Scheduler metrics suspicious: high initial latency (60ms),  
deteriorating over time (~200ms)

Scheduling latency increasing with total number of scheduled pods

Isolated to scheduler, but still in a complex codebase...

Fine-grained profiling required!

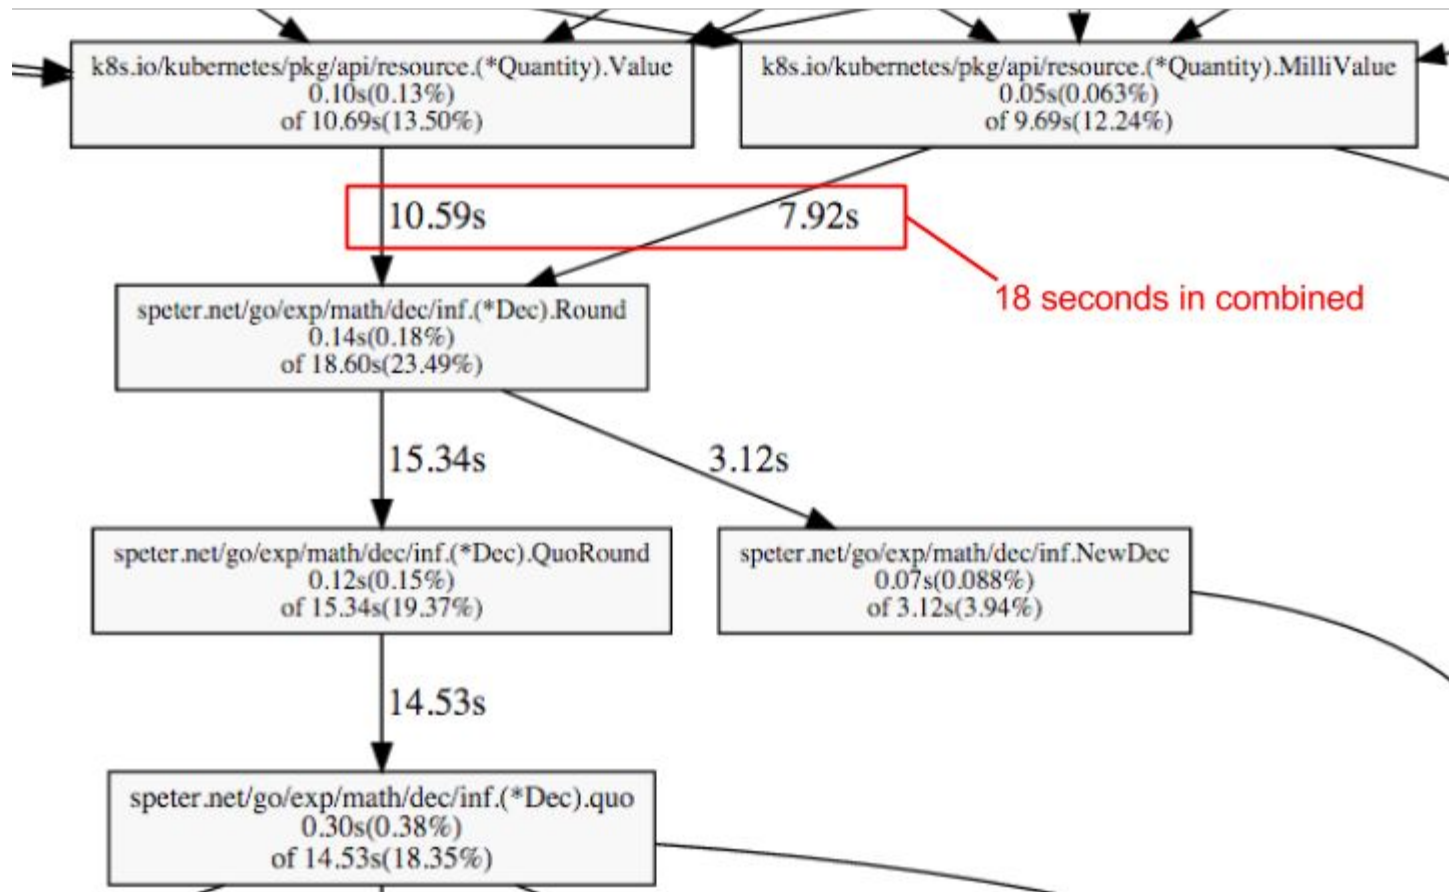
# Taking a shortcut

Kubemark too slow: hours to run each test

Wrote a lightweight [custom benchmarking tool](#)

Analysed the go runtime performance with pprof

File: benchmark.test  
Type: cpu  
54.09s of 79.19s total (68.30%)



# Taking a shortcut

Several issues identified, fixes proposed:

- <https://github.com/kubernetes/kubernetes/issues/18126>
- <https://github.com/kubernetes/kubernetes/pull/18170>
- <https://github.com/kubernetes/kubernetes/issues/18255>
- <https://github.com/kubernetes/kubernetes/pull/18413>
- <https://github.com/kubernetes/kubernetes/issues/18831>

## Case study: end results

- Average scheduling latency reduced from **53 seconds** to **23 seconds** for 1 000 pods on 1 000 nodes
- Time to schedule 30 000 pods onto 1 000 nodes reduced from **8 780** seconds to **587** seconds
- 4 merged Kubernetes pull requests
- Prometheus, top, logplot, pprof, and more...

# **Case study: what went wrong?**

Taking full(er) advantage of metrics and Prometheus

# Case study: what went wrong?

Using 3 cores out of 16 cores

PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
20	0	268m	253m	23m	S	192	0.4	4:25.43	kube-apiserver
20	0	58000	47m	15m	S	53	0.1	0:51.18	kube-scheduler
20	0	130m	118m	20m	S	17	0.2	0:27.61	kube-controller
20	0	124m	111m	10m	S	15	0.2	0:21.31	etcd

*Resource usage during the test*

# Case study: what went wrong?

Using 3 cores out of 16 cores

PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
20	0	268m	253m	23m	S	192	0.4	4:25.43	kube-apiserver
20	0	58000	47m	15m	S	33	0.1	0:51.18	kube-scheduler
20	0	130m	118m	20m	S	17	0.2	0:27.61	kube-controller
20	0	124m	111m	10m	S	15	0.2	0:21.31	etcd

*Resource usage during the test*

**Not** a good example of metrics debugging!



# Case study: what went wrong?

Using 3 cores out of 16 cores

PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
20	0	268m	253m	23m	S	192	0.4	4:25.43	kube-apiserver
20	0	58000	47m	15m	S	33	0.1	0:51.18	kube-scheduler
20	0	130m	118m	20m	S	17	0.2	0:27.61	kube-controller
20	0	124m	111m	10m	S	15	0.2	0:21.31	etcd

*Resource usage during the test*

**Better:** use Prometheus for system level metrics


# Case study: what went wrong?

**Kubemark** test used log scraping + custom graphing tool

```
Pods: 229 out of 3000 created, 211 running, 18 pending, 0 waiting
Pods: 429 out of 3000 created, 412 running, 17 pending, 0 waiting
Pods: 622 out of 3000 created, 604 running, 17 pending, 1 waiting
...
Pods: 2578 out of 3000 created, 2561 running, 17 pending, 0 waiting
Pods: 2779 out of 3000 created, 2761 running, 18 pending, 0 waiting
Pods: 2979 out of 3000 created, 2962 running, 16 pending, 1 waiting
Pods: 3000 out of 3000 created, 3000 running, 0 pending, 0 waiting
```

# Case study: what went wrong?

**Kubemark** test used log scraping + custom graphing tool



```
Pods: 229 out of 3000 created, 211 running, 18 pending, 0 waiting
Pods: 429 out of 3000 created, 412 running, 17 pending, 0 waiting
Pods: 622 out of 3000 created, 604 running, 17 pending, 1 waiting
...
Pods: 2578 out of 3000 created, 2561 running, 17 pending, 0 waiting
Pods: 2779 out of 3000 created, 2761 running, 18 pending, 0 waiting
Pods: 2979 out of 3000 created, 2962 running, 16 pending, 1 waiting
Pods: 3000 out of 3000 created, 3000 running, 0 pending, 0 waiting
```

**Better:** use Prometheus + Grafana (or similar)

# Why are metrics better?

Simplify the lower level analysis

- Access to meaningful internal information without spelunking through code
- No need to reproduce the problem
  - No need to attach debugger
- Dynamic power of a query language
  - No static SVG graph (pprof)


# Why are metrics better?

More powerful higher level analysis

- Fleet-wide view and aggregation
- Historic view
- Easy correlation with other internal and external factors over time

tl;dr

- Metrics don't provide all the answers (e.g. profiling a single function call), but in any complex system they help understand and debug behaviour and guide further investigation.



Core OS <3



# Join us!

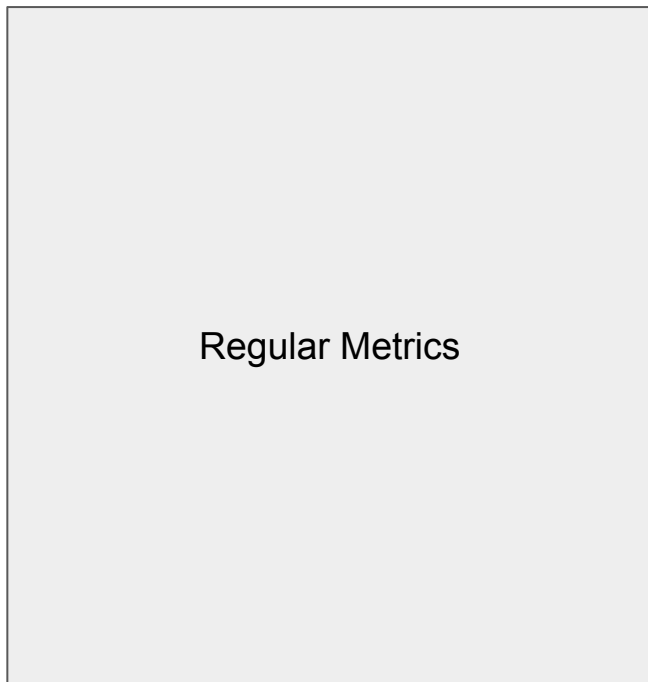
Hiring Prometheus developers in Berlin

[coreos.com/careers](https://coreos.com/careers)





`/metrics`



Regular Metrics

`/debug_metrics`



Debugging Metrics

ingest optionally

`/metrics`

Regular Metrics

Debugging Metrics

ingest by default,  
ignore optionally