

The Open Container Initiative

Establishing standards for an open ecosystem

Jonathan Boule

@baronboule | jon@coreos.com

~~The Open Container Initiative~~

~~Establishing standards for an open ecosystem~~

Containers, standards, and the pianola

In search of a better metaphor

Jonathan Boule

@baronboule | jon@coreos.com

~~The Open Container Initiative~~

~~Establishing standards for an open ecosystem~~

Containers, standards, and the pianola

In search of a better metaphor

(...with some stuff about OCI too)

Jonathan Boule

@baronboule | jon@coreos.com

Containers, standards, and the pianola

... a tale of

- shipping containers
- infinite software
- cheesecake
- IKEA furniture
- mechanical pianos (aka the pianola)
- container standards

But first, a bit about OCI

OPEN CONTAINER INITIATIVE

AN OPEN GOVERNANCE STRUCTURE FOR THE
EXPRESS PURPOSE OF CREATING OPEN INDUSTRY
STANDARDS AROUND CONTAINER FORMATS AND
RUNTIME

Containers and OCI

- OCI - the Open Container Initiative
 - Founded mid-2015 to end the "Container Wars"

Containers and OCI

- OCI - the Open Container Initiative
 - Founded mid-2015 to end the "Container Wars"
 - But really... to formalise the *de facto* Docker standard
 - Agreed-on starting point for future innovation

Containers and OCI

- OCI - the Open Container Initiative
 - Founded mid-2015 to end the "Container Wars"
 - But really... to formalise the *de facto* Docker standard
 - Agreed-on starting point for future innovation
- Two key projects
 - *Image specification* - what's in a container
 - *Runtime specification* - how a container runs

Containers and OCI

- What even is an application container?
 - cgroups? Namespaces? Union filesystems?

Containers and OCI

- What even is an application container?
 - cgroups? Namespaces? Union filesystems?
- Answer: not very interesting
 - A tarball and a bunch of JSON metadata

Containers and metaphors

Containers and (clichéd) metaphors



Shipping containers

- ✓ Agreed-on format (size and shape)
- ✓ Works with cranes, ships, trucks, trains, ...
- ✓ Transports can ignore what's inside
- ✓ Consistent experience

Application containers

- ✓ Agreed-on format
- ✓ Works with registries, build tools, runtimes...
- ✓ Transports can ignore what's inside
- ✓ Consistent experience



Shipping containers

- ✗ Operators can ignore what's inside
 - Because it's opaque and unstructured

What's inside containers?

- Shipping containers
 - Yoghurt?
 - Furniture?
 - Smaller containers? (containers all the way down)
 - Don't know how to start dealing with the contents

What's inside containers?


- Shipping containers
 - Yoghurt?
 - Furniture?
 - Smaller containers? (containers all the way down)
 - Don't know how to start dealing with the contents
- Application containers
 - A structured filesystem layout
 - An entrypoint: start with `/bin/httpd`

Shipping containers

- ✗ Operators can ignore what's inside
 - Because it's opaque and unstructured

Shipping containers

- ✓ Operators can ignore what's inside
 - Add a shipping manifest!
 - What's inside, how to process it



Shipping containers

× Monolithic size

- Each container holds the same amount
- Application containers can vary wildly

Shipping containers

- ✗ Monolithic size
 - Each container holds the same amount
 - Application containers can vary wildly
- ✗ Physically cumbersome
 - Difficult to build, difficult to move
 - Application containers can be copied in an instant

Shipping containers

- ✗ Monolithic size
 - Each container holds the same amount
 - Application containers can vary wildly
- ✗ Physically cumbersome
 - Difficult to build, difficult to move
 - Application containers can be copied in an instant
- ✗ Filled, emptied, re-used
 - Application containers are immutable, copied

Physical metaphors are hard


Software is

- instantly, immediately, cheaply copied
- instantly, immediately, cheaply transported
- only constrained by supply of electricity

Physical metaphors are hard

Software is

- a stream of bits: ones and zeroes
 - true, but not very helpful



Physical metaphors are hard

Software is

- ~~a stream of bits: ones and zeroes~~
 - ~~true, but not very helpful~~
- a sequence of instructions, potentially endless
 - CPU dumbly follows these instructions (but really fast)
 - recreate the sequence, recreate the software
 - copy the sequence, copy the software


Physical metaphors - trying again

Sequence of instructions?

What about a recipe?

Recipes

- ✓ Sequence of instructions
- ✓ Easy to copy and redistribute
- ✓ Follow the instructions, get the same result



Recipes

- ✓ Sequence of instructions
- ✓ Easy to copy and redistribute
- ✓ Follow the instructions, get the same result



Application containers

- ✓ Sequence of instructions
- ✓ Easy to copy and redistribute
- ✓ Follow the instructions, get the same result

Application containers

- ✓ Sequence of instructions
- ✓ Easy to copy and redistribute
- ✓ Follow the instructions, get the same result



Recipes



- ✗ Needs an additional set of inputs (ingredients)
 - Inconsistency in results
 - Not self-contained like an application container

Recipes

- ✗ Needs an additional set of inputs (ingredients)
 - Inconsistency in results
 - Not self-contained like an application container



IKEA furniture



IKEA furniture

- ✓ Sequence of instructions
- ✓ Easy to copy and redistribute
- ✓ Follow the instructions, get the same result
- ✓ Ingredients (materials) included in the package
 - Self-contained, consistent result!

IKEA furniture



Physical metaphors - improving!

Physical metaphors are hard

Software is

- instantly, immediately, cheaply copied
- instantly, immediately, cheaply transported
- only constrained by supply of electricity

Physical metaphors are hard

Software is

- instantly, immediately, cheaply copied
- instantly, immediately, cheaply transported
- only constrained by supply of electricity
- **As long as you have electricity, software is *long-running* and *dynamic* (alive)**

Things that are not alive

- ✗ Shipping containers are (relatively) static
- ✗ So are IKEA bookshelves
- ✗ Cheesecakes get eaten
 - If not, let me know



slackbot 14:30

Reminder: get cheesecake!

Physical metaphors are hard

- ✗ Shipping containers are (relatively) static
 - ✗ So are IKEA bookshelves
 - ✗ Cheesecakes get eaten
-
- ✓ Software goes on,
and on,
and on



Physical metaphors - one more try






Containers and the pianola

- **Piano roll as application containers**
 - Set of instructions to follow
- **Pianola as computer** (specifically, CPU)
 - Dumbly follows instructions, only input is energy
- **Musical performance as software execution**
 - Dynamic, alive, long-lived

The piano roll

Packaging format

- Machine-understandable sheet music
- Various lengths, sizes
- Easy to copy, transport
- Can reference each other



The pianola

- Mechanical piano
- Dumbly consumes piano roll
- Different energy sources (e.g. pneumatic)



Software is alive

- Long running, dynamic, self-referential
- Limitless (with enough electricity)

Software is alive

- Long running, dynamic, self-referential
- Limitless (with enough electricity)



So is musical performance!

- Repeat a code block, repeat a bar of music
- Limitless (with enough mechanical energy)

How far does this go?

Sheet music and container layers

- Software containers have *layers*
 - One application container derives from another
 - No need to copy all of the files you need; simply reference the parent layer

Sheet music and container layers

- Software containers have *layers*
 - One application container derives from another
 - No need to copy all of the files you need; simply reference the parent layer
- Piano rolls can, too!
 - Reference a bar or piece from another roll
 - Share common melodies, etc

Sheet music and container standards

- Standardise *entrypoint*

Sheet music and container standards

- Standardise *entrypoint*
 - How can the pianola tell where in the sheet music it should start playing?
 - "entrypoint": "19,2"

Sheet music and container standards

- Standardise *entrypoint*
 - How can the pianola tell where in the sheet music it should start playing?
 - "entrypoint": "19,23"
 - How can the container runtime tell which binary it should start executing?
 - "entrypoint": "/bin/httpd"

Sheet music and container standards

- Standardise *constraints*

Sheet music and container standards

- Standardise *constraints*
 - How loud can this piece be played?
 - "maxVolume": "123dB"
 - Exceed the limit? Music stops

Sheet music and container standards

- Standardise *constraints*
 - How loud can this piece be played?
 - "maxVolume": "123dB"
 - Exceed the limit? Music stops
 - How much memory can this container use?
 - "maxMemory": "123MB"
 - Exceed the limit? Software stops

Sheet music and container standards

- Standardise *discovery*

Sheet music and container standards

- Standardise *discovery*
 - How can I find this referenced piece by Bono?
 - Look up Bono in the telephone book
 - Call the phone number
 - Ask for his cool piece

Sheet music and container standards

- Standardise *discovery*
 - How can I find this referenced piece by Bono?
 - Look up Bono in the telephone book
 - Call the phone number
 - Ask for his cool piece
 - How can I find this referenced container image layer?
 - Look up bono.com in DNS
 - Connect to port 80
 - GET /songs/cool_piece JPTP/1.1

What else?

Multiple clients / listeners

Remote access / listening

Container orchestration / pianola orchestras

So much more...

Time check-in

And maybe some more on OCI

OCI Today

Two separate but connected specifications

- **image-spec**: what's in a container
- **runtime-spec**: how to run a container

OCI Image Format Spec Project

- A serialized, distributable image format
 - Content-addressable
 - Platform-agnostic
- Optional extras:
 - Signatures based on image content address
 - Federated, delegatable naming based on DNS


OCI Image Format Spec Project

- Portable archive format
 - "The thing to distribute"
 - Structured tarball
- Image Manifest, Image Index, and Config
 - Metadata about the container image
 - References to *layers*, containing root filesystem
- Cryptographic addressability
 - How to securely reference images and parts of images

OCI Image Format Spec Project

- Backwards-compatible with Docker:
 - Taking the *de facto* standard Docker v2.2 format and writing it down for everyone to use
- Shared starting point for future innovation in container image format and distribution
- Intended to interoperate with Runtime Spec (similar to how appc defined both sections)

Anatomy of an OCI Image



Inside the tarball

```
$ find busybox/  
busybox/  
busybox/refs  
busybox/refs/latest  
busybox/oci-layout  
busybox/blobs  
busybox/blobs/sha256  
busybox/blobs/sha256/d09bddf0432...  
busybox/blobs/sha256/56bec22e355...  
busybox/blobs/sha256/e02e811dd08...
```

```
$ cat busybox/blobs/sha256/d09bddf043...  
{  
  "layers" : [  
    { "digest" : "sha256:56bec22e355981d...",  
      "size" : 668151,  
      "mediaType" : application/vnd.oci.image.layer.v1.tar+gzip"  
    } ],  
  "mediaType" : "application/vnd.oci.image.manifest.v1+json",  
  "schemaVersion" : 2,  
  "config" : {  
    "digest" : "sha256:e02e811dd08fd49e7f6...",  
    "mediaType" : "application/vnd.oci.image.config.v1+json",  
    "size" : 1464  
  }  
}
```

OCI Runtime Spec

- On-disk layout of a container
 - Extracted root filesystem and configuration, ready to run
- Lifecycle verbs
 - create, start, kill, delete, state
- Multi-platform support
 - Shared general configuration
 - Windows/Solaris/Linux-specific bits

OCI Runtime Spec

Example: container state

```
{  
  "ociVersion": "v1.0.0-rc5",  
  "id": "oci-container1",  
  "status": "running",  
  "pid": 4422,  
  "bundlePath": "/containers/redis",  
  "annotations": {  
    "myKey": "myValue"  
  }  
}
```

Thank you!

All OCI work happens in the open - join us!

- GitHub:
 - <https://github.com/opencontainers/image-spec>
 - <https://github.com/opencontainers/runtime-spec>
- Email:
 - dev@opencontainers.org

TRUMP SEEKS "THE ULTIMATE DEAL"

...HOW MUCH
FOR THE WALL?

DAILY MAVERICK
23-5-17
ZAPERO

